

eda_preprocessing

February 21, 2026

1 EDA and Preprocessing

This notebook covers the Exploratory Data Analysis (EDA) and preprocessing steps for the AAI-530 final project.

Goals:

- Load and validate the AI4I 2020 Predictive Maintenance dataset
- Perform core EDA (types, missing values, duplicates, target distribution)
- Explore feature distributions, feature-to-target relationships, and correlations
- Document the project-specific assumption of using `Tool wear [min]` as a proxy for time progression
- Prepare a leakage-aware modeling dataset (drop IDs, remove failure-mode flags, encode categorical features)
- Export a prepared dataset for downstream modeling notebooks

1.0.1 1. Imports & setup

```
[1]: import pandas as pd
      import numpy as np

      import matplotlib.pyplot as plt
      import seaborn as sns

      # Display settings
      pd.set_option("display.max_columns", None)
      sns.set(style="whitegrid")
```

1.0.2 2. Load dataset and quick validation

```
[2]: # Path to dataset
DATA_PATH = "../data/ai4i_2020_predictive_maintenance.csv"

# Load CSV
df = pd.read_csv(DATA_PATH)

# Basic check
df.head()
```

```
[2]:   UDI Product ID Type Air temperature [K] Process temperature [K] \
0     1      M14860     M             298.1            308.6
1     2      L47181     L             298.2            308.7
```

2	3	L47182	L	298.1	308.5		
3	4	L47183	L	298.2	308.6		
4	5	L47184	L	298.2	308.7		
0		Rotational speed [rpm]	Torque [Nm]	Tool wear [min]	Machine failure	TWF	\
0		1551	42.8	0	0	0	0
1		1408	46.3	3	0	0	0
2		1498	49.4	5	0	0	0
3		1433	39.5	7	0	0	0
4		1408	40.0	9	0	0	0
0	HDF	PWF	OSF	RNF			
0	0	0	0	0			
1	0	0	0	0			
2	0	0	0	0			
3	0	0	0	0			
4	0	0	0	0			

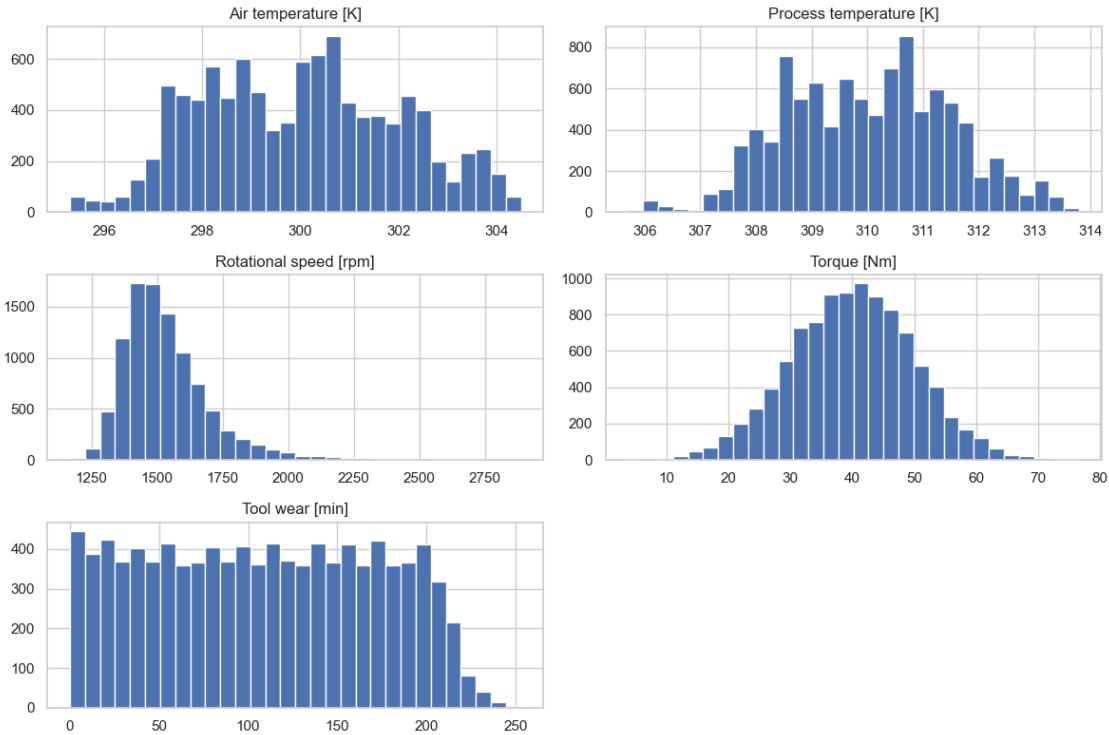
1.0.3 3. Feature distributions and outliers

This section checks basic feature distributions to highlight skewness and potential outliers before modeling.

```
[3]: import matplotlib.pyplot as plt

numeric_cols = [
    'Air temperature [K]',
    'Process temperature [K]',
    'Rotational speed [rpm]',
    'Torque [Nm]',
    'Tool wear [min]'
]

df[numeric_cols].hist(bins=30, figsize=(12, 8))
plt.tight_layout()
plt.show()
```



1.0.4 4. Dataset structure and summary statistics

```
[4]: print("Dataset shape:", df.shape)
print("\nColumn info:")
df.info()
```

Dataset shape: (10000, 14)

Column info:

```
<class 'pandas.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   UDI              10000 non-null   int64  
 1   Product ID       10000 non-null   str    
 2   Type              10000 non-null   str    
 3   Air temperature [K] 10000 non-null   float64
 4   Process temperature [K] 10000 non-null   float64
 5   Rotational speed [rpm] 10000 non-null   int64  
 6   Torque [Nm]        10000 non-null   float64
 7   Tool wear [min]    10000 non-null   int64  
 8   Machine failure    10000 non-null   int64  
 9   TWF               10000 non-null   int64
```

```

10  HDF          10000 non-null  int64
11  PWF          10000 non-null  int64
12  OSF          10000 non-null  int64
13  RNF          10000 non-null  int64
dtypes: float64(3), int64(9), str(2)
memory usage: 1.1 MB

```

Summary statistics

```
[5]: df.describe(include="all")
```

	UDI	Product ID	Type	Air temperature [K]	\
count	10000.00000	10000	10000	10000.000000	
unique	Nan	10000	3	Nan	
top	Nan	M14860	L	Nan	
freq	Nan	1	6000	Nan	
mean	5000.50000	Nan	Nan	300.004930	
std	2886.89568	Nan	Nan	2.000259	
min	1.00000	Nan	Nan	295.300000	
25%	2500.75000	Nan	Nan	298.300000	
50%	5000.50000	Nan	Nan	300.100000	
75%	7500.25000	Nan	Nan	301.500000	
max	10000.00000	Nan	Nan	304.500000	
	Process temperature [K]	Rotational speed [rpm]	Torque [Nm]	\	
count	10000.000000	10000.000000	10000.000000		
unique	Nan	Nan	Nan	Nan	
top	Nan	Nan	Nan	Nan	
freq	Nan	Nan	Nan	Nan	
mean	310.005560	1538.776100	39.986910		
std	1.483734	179.284096	9.968934		
min	305.700000	1168.000000	3.800000		
25%	308.800000	1423.000000	33.200000		
50%	310.100000	1503.000000	40.100000		
75%	311.100000	1612.000000	46.800000		
max	313.800000	2886.000000	76.600000		
	Tool wear [min]	Machine failure	TWF	HDF	\
count	10000.000000	10000.000000	10000.000000	10000.000000	
unique	Nan	Nan	Nan	Nan	
top	Nan	Nan	Nan	Nan	
freq	Nan	Nan	Nan	Nan	
mean	107.951000	0.033900	0.004600	0.011500	
std	63.654147	0.180981	0.067671	0.106625	
min	0.000000	0.000000	0.000000	0.000000	
25%	53.000000	0.000000	0.000000	0.000000	
50%	108.000000	0.000000	0.000000	0.000000	
75%	162.000000	0.000000	0.000000	0.000000	

max	253.000000	1.000000	1.000000	1.000000
	PWF	OSF	RNF	
count	10000.000000	10000.000000	10000.000000	
unique	NaN	NaN	NaN	
top	NaN	NaN	NaN	
freq	NaN	NaN	NaN	
mean	0.009500	0.009800	0.001900	
std	0.097009	0.098514	0.043550	
min	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.000000	
max	1.000000	1.000000	1.000000	

1.0.5 5. Data quality checks (missing, uniqueness, duplicates)

```
[6]: missing= df.isnull().sum()
missing[missing > 0]
```

```
[6]: Series([], dtype: int64)
```

Check uniqueness and duplicates

```
[7]: df.nunique().sort_values()
df.duplicated().sum()

print(f"Duplicate rows: {df.duplicated().sum()}")
```

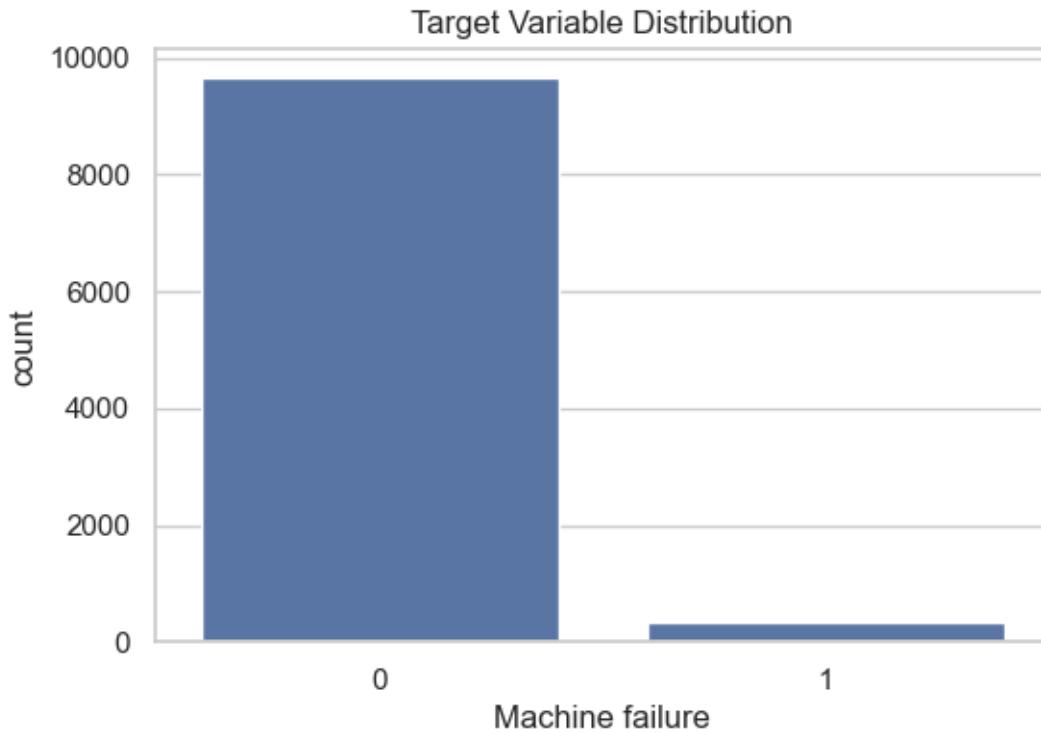
Duplicate rows: 0

1.0.6 6. Target variable analysis

```
[8]: # Assuming 'Machine failure' is the target
target_col = "Machine failure"

df[target_col].value_counts(normalize=True)

plt.figure(figsize=(6, 4))
sns.countplot(x=target_col, data=df)
plt.title("Target Variable Distribution")
plt.show()
```



The target variable is highly imbalanced, with failures representing a small fraction of observations, which motivates the use of stratified sampling and appropriate evaluation metrics in downstream models.

1.0.7 7. Identifying the time-series variable

```
[9]: df['Tool wear [min]'].describe()
```

```
[9]: count      10000.000000
mean       107.951000
std        63.654147
min        0.000000
25%       53.000000
50%      108.000000
75%      162.000000
max      253.000000
Name: Tool wear [min], dtype: float64
```

1.0.8 8. Time proxy exploration using tool wear

Assumption: using tool wear as a proxy for time

The AI4I dataset is not a single-machine chronological log. For this project, we treat **Tool wear [min]** as a proxy for progression over time-under-use. When we aggregate and visualize trends by

tool wear, results reflect population-level patterns rather than an individual machine trajectory.

```
[10]: df = df.sort_values('Tool wear [min]').reset_index(drop=True)
```

1.0.9 Failure vs key features

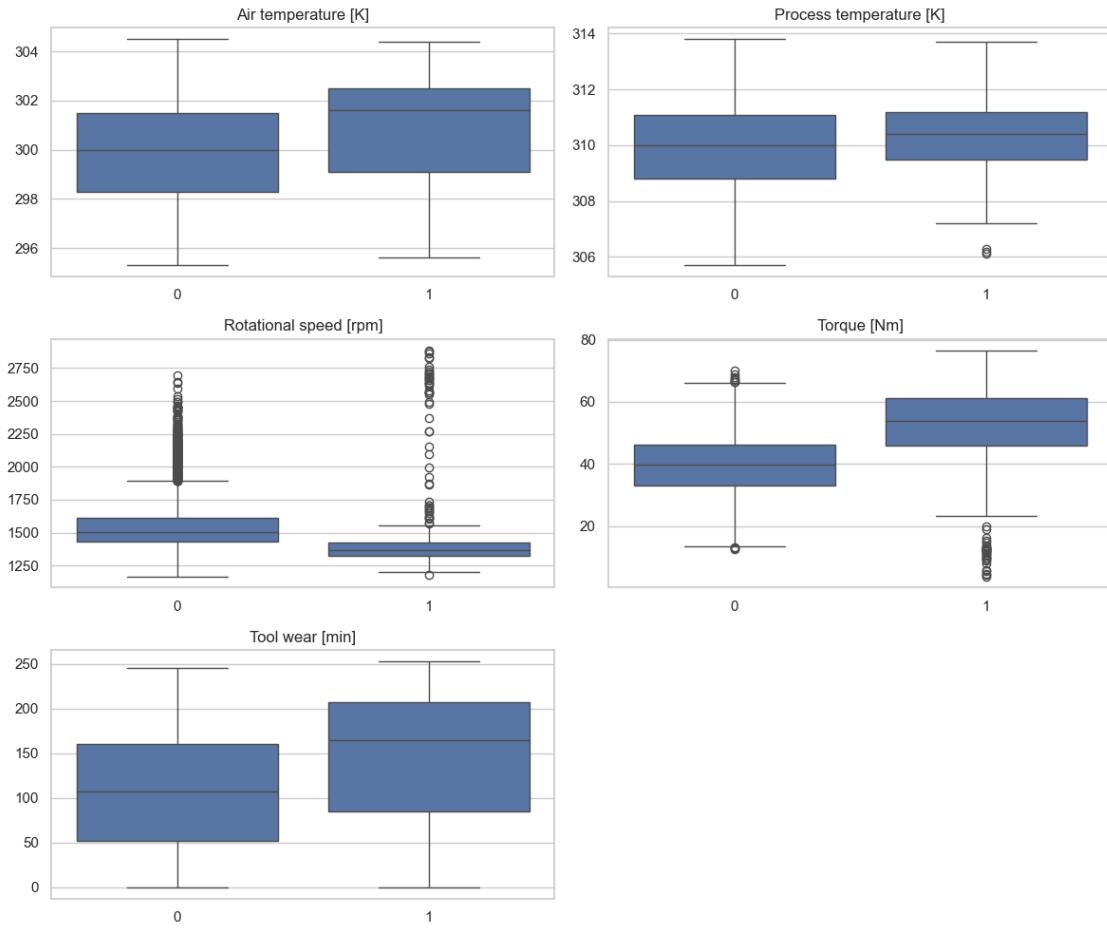
```
[11]: key_features = [
    "Air temperature [K]",
    "Process temperature [K]",
    "Rotational speed [rpm]",
    "Torque [Nm]",
    "Tool wear [min]"
]

fig, axes = plt.subplots(3, 2, figsize=(12, 10))
axes = axes.flatten()

for ax, feature in zip(axes, key_features):
    sns.boxplot(x=target_col, y=feature, data=df, ax=ax)
    ax.set_title(feature)
    ax.set_xlabel("")
    ax.set_ylabel("")

for ax in axes[len(key_features):]:
    ax.remove()

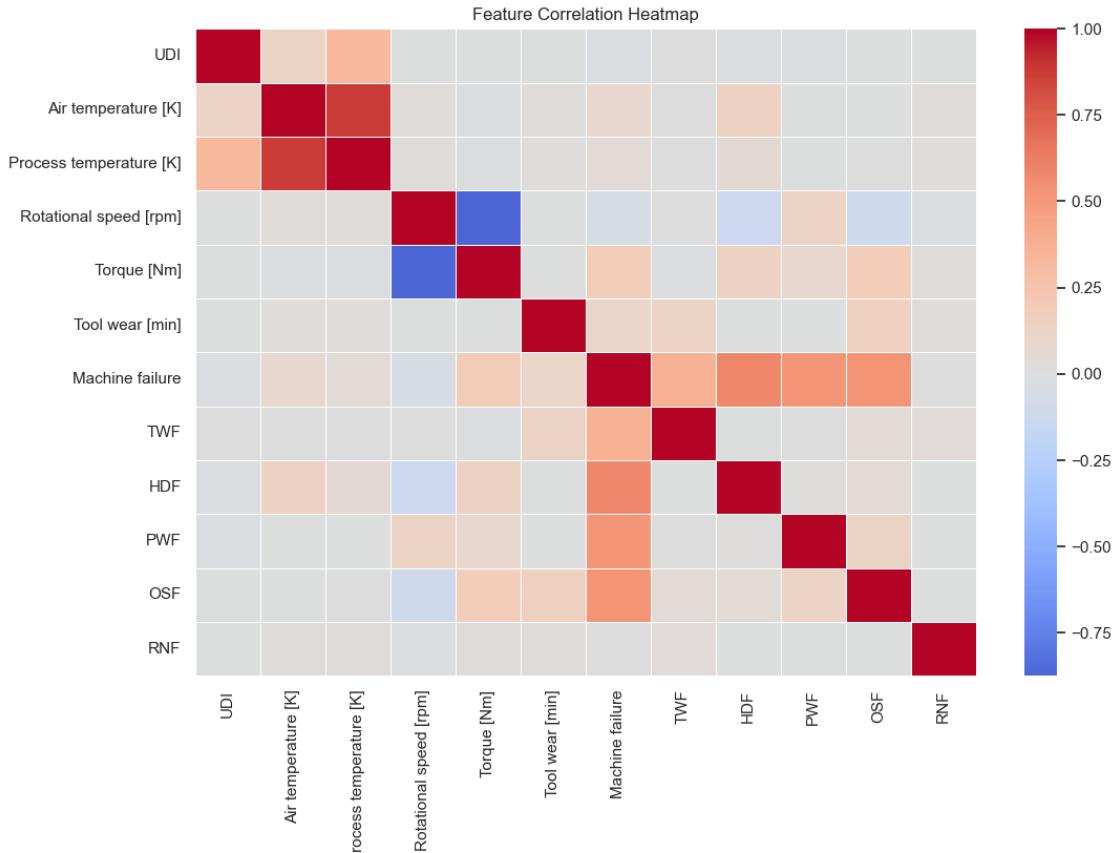
plt.tight_layout()
plt.show()
```



1.0.10 9. Correlation analysis

```
[12]: numerical_cols = df.select_dtypes(include=np.number).columns.tolist()
plt.figure(figsize=(12, 8))
corr = df[numerical_cols].corr()

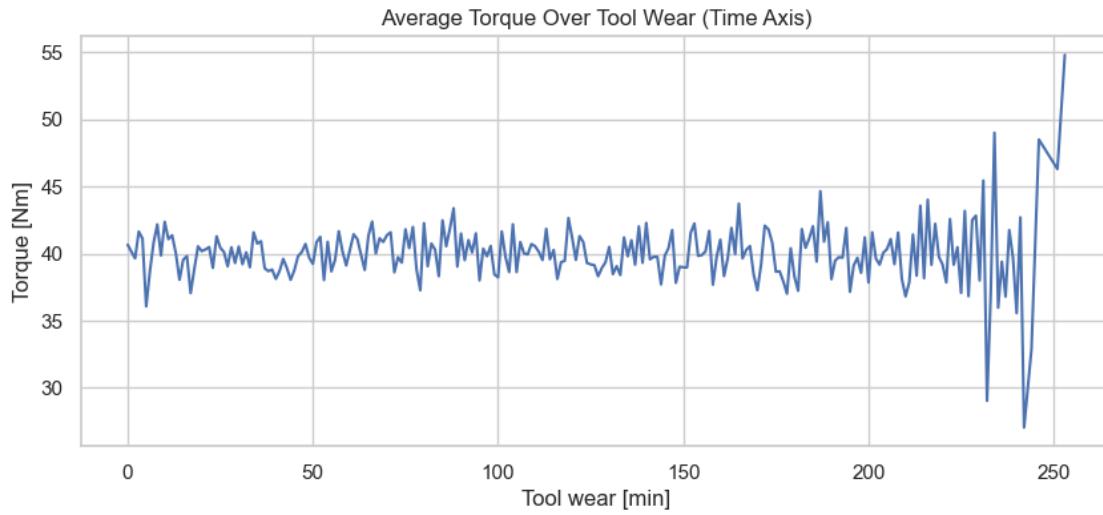
sns.heatmap(
    corr,
    cmap="coolwarm",
    center=0,
    linewidths=0.5
)
plt.title("Feature Correlation Heatmap")
plt.show()
```



1.0.11 10. Time-series behavior using Tool wear

```
[13]: df_tw = df.groupby('Tool wear [min]', as_index=False).mean(numeric_only=True)

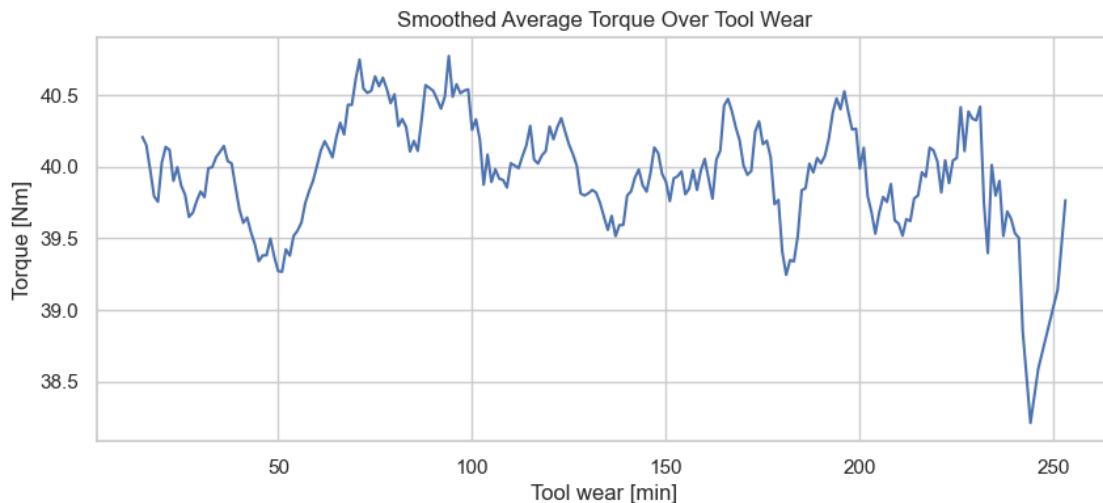
plt.figure(figsize=(10,4))
plt.plot(df_tw['Tool wear [min]'], df_tw['Torque [Nm]'])
plt.title("Average Torque Over Tool Wear (Time Axis)")
plt.xlabel("Tool wear [min]")
plt.ylabel("Torque [Nm]")
plt.show()
```



1.0.12 11. Moving average smoothing to reduce noise

```
[14]: df_tw['Torque_ma'] = df_tw['Torque [Nm]'].rolling(window=15).mean()
```

```
plt.figure(figsize=(10,4))
plt.plot(df_tw['Tool wear [min]', df_tw['Torque_ma']])
plt.title("Smoothed Average Torque Over Tool Wear")
plt.xlabel("Tool wear [min]")
plt.ylabel("Torque [Nm]")
plt.show()
```



1.0.13 12. Preprocessing for modeling (leakage-aware)

```
[15]: # Prepare dataset for ML prep steps (remove IDs and leakage variables)
drop_cols = ['UDI', 'Product ID', 'TWF', 'HDF', 'PWF', 'OSF', 'RNF']
df_prepared = df.drop(columns=drop_cols)

print("Dropped columns:", drop_cols)
print("Prepared dataset shape:", df_prepared.shape)
```

Dropped columns: ['UDI', 'Product ID', 'TWF', 'HDF', 'PWF', 'OSF', 'RNF']
Prepared dataset shape: (10000, 7)

Identifier fields and failure-mode indicators were removed prior to modeling to prevent data leakage and ensure that predictions are based only on sensor measurements and operational variables.

```
[16]: #Encoding
df_prepared = pd.get_dummies(df_prepared, columns=['Type'], drop_first=True)
```

1.0.14 13. Export prepared dataset

```
[17]: # Save prepared dataset for modeling notebooks (optional)
OUTPUT_PATH = "../data/ai4i_prepared.csv"
df.to_csv(OUTPUT_PATH, index=False)
print(f"Saved prepared dataset to: {OUTPUT_PATH}")
```

Saved prepared dataset to: ../data/ai4i_prepared.csv

1.0.15 14. EDA Summary

1. Data loading and overview

- Load the dataset and display sample rows
- Verify dataset shape, column names, and basic structure

2. Data types and summary statistics

- Review data types using `info()`
- Review basic statistics using `describe()`

3. Data quality checks

- Check missing values and confirm handling approach
- Check unique values and duplicate rows

4. Target variable analysis

- Examine the distribution of the target (`Machine failure`)
- Note class imbalance and modeling implications

5. Feature exploration

- Inspect distributions of key numeric features
- Compare feature behavior between failure and non-failure cases

6. Correlation analysis

- Compute and visualize correlations among numeric features
- Identify highly correlated variables

7. Time proxy assumption

- Justify using `Tool wear [min]` as a proxy for time
- Visualize aggregate trends across tool wear
- State limitations of this assumption

8. Preprocessing for modeling

- Remove non-predictive identifiers
- Drop leakage-prone failure mode flags
- Encode categorical features
- Save the prepared dataset for downstream modeling

model1A_classification_machine_failure

February 21, 2026

1 Model 1A — Traditional ML — Time-Aware Failure Classification

This notebook trains a traditional machine learning model using **time-aware features** derived from sensor readings.

Project context: - Dataset: AI4I 2020 Predictive Maintenance (industrial sensor snapshots) - Time proxy: Tool wear [min] is treated as progression over time-under-use - Goal (Model 1): Predict **failure risk** as tool wear increases using traditional ML (baseline + tree-based)

Notes: - We split train/test by **time order** (no shuffle). - We avoid label leakage by excluding the failure-mode flags when predicting **Machine failure**.

1.0.1 1. Import & Setup

```
[1]: import os
import warnings
warnings.filterwarnings("ignore")

import numpy as np
import pandas as pd

from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.metrics import (
    classification_report,
    confusion_matrix,
    roc_auc_score,
    roc_curve,
    average_precision_score,
    RocCurveDisplay,
    PrecisionRecallDisplay
)
import matplotlib.pyplot as plt
```

```
pd.set_option("display.max_columns", 200)
pd.set_option("display.width", 120)
```

1.0.2 2. Load prepared data

We prefer using the exported prepared dataset from the EDA notebook (`data/ai4i_prepared.csv`). If it is not available, we load the raw dataset and apply minimal preprocessing:

- Drop identifiers (UDI, Product ID)
- Drop failure-mode flags (TWF, HDF, PWF, OSF, RNF) to avoid leakage when predicting Machine failure
- One-hot encode Type

```
[2]: DATA_PREPARED_PATH = "../data/ai4i_prepared.csv"
DATA_RAW_PATH = "../data/ai4i_2020_predictive_maintenance.csv"

TARGET_COL = "Machine failure"
TIME_COL = "Tool wear [min]"

def load_dataset():
    if os.path.exists(DATA_PREPARED_PATH):
        df = pd.read_csv(DATA_PREPARED_PATH)
        source = "prepared"
    else:
        df_raw = pd.read_csv(DATA_RAW_PATH)
        drop_cols = ["UDI", "Product ID", "TWF", "HDF", "PWF", "OSF", "RNF"]
        df = df_raw.drop(columns=[c for c in drop_cols if c in df_raw.columns], u
        ↪errors="ignore")
        if "Type" in df.columns:
            df = pd.get_dummies(df, columns=["Type"], drop_first=True)
            source = "raw+prepped"
    return df, source

df, source = load_dataset()
print(f"Loaded source: {source}")
print("Shape:", df.shape)
df.head()
```

Loaded source: prepared
 Shape: (10000, 14)

	UDI	Product ID	Type	Air temperature [K]	Process temperature [K]
	Rotational speed [rpm]			Torque [Nm]	\
0	1	M14860	M	298.1	308.6
1551			42.8		
1	7257	H36670	H	300.2	310.3
1408			42.5		
2	504	M15363	M	297.6	309.2
1442			48.1		
3	7169	L54348	L	300.3	310.3
1704			29.5		

4	7089	M21948	M	300.6	310.3
1614			32.7		
		Tool wear [min]	Machine failure	TWF	HDF
0		0	0	0	PWF
1		0	0	0	OSF
2		0	0	0	RNF
3		0	0	0	
4		0	0	0	

1.0.3 3. Validate required columns

```
[3]: sensor_cols = [
    "Air temperature [K]",
    "Process temperature [K]",
    "Rotational speed [rpm]",
    "Torque [Nm]",
    TIME_COL,
    TARGET_COL
]

missing = [c for c in sensor_cols if c not in df.columns]
if missing:
    raise ValueError(f"Missing required columns: {missing}")

print("All required columns are present.")
```

All required columns are present.

1.0.4 4. Time-aware feature engineering

Because AI4I is not a per-machine chronological log, we use a **population-level** time proxy:

- Sort by **Tool wear [min]** - Aggregate to a single series over tool wear (mean of numeric features)
- Create lag, rolling mean, rolling std, and deltas to capture change over time

This creates a clean, ordered sequence suitable for traditional time-series style modeling.

```
[4]: df_tw = (
    df
    .groupby(TIME_COL, as_index=False)
    .agg({
        "Air temperature [K]": "mean",
        "Process temperature [K]": "mean",
        "Rotational speed [rpm)": "mean",
        "Torque [Nm]": "mean",
        TARGET_COL: "max",
    })
    .sort_values(TIME_COL)
```

```

    .reset_index(drop=True)
)

print("Aggregated shape:", df_tw.shape)
print("Label distribution after aggregation:")
print(df_tw[TARGET_COL].value_counts())
df_tw.head()

```

Aggregated shape: (246, 6)
Label distribution after aggregation:
Machine failure
1 172
0 74
Name: count, dtype: int64

[4]:

	Tool wear [min]	Air temperature [K]	Process temperature [K]	Rotational speed [rpm]	Torque [Nm]	Machine failure
0		299.956667			0	309.955833
1524.916667	40.661667		1		2	300.272464
1					3	310.142029
1555.521739	39.646377		1		4	299.679412
2					5	309.826471
1508.264706	41.644118		1		6	299.997059
3					7	309.870588
1525.882353	41.117647		0		8	299.925397
4					9	310.014286
1620.761905	36.071429		1		10	

1.0.5 5. Create time-aware features (lags, rolling stats, deltas)

We create features that depend only on past values:

- Lag features
- Rolling mean/std (past-only)
- First differences (deltas)

[5]:

```

def make_time_features(
    df_time: pd.DataFrame,
    base_cols,
    time_col: str,
    target_col: str,
    windows=(3, 5, 10),
    lags=(1, 2, 3),
):
    # Avoid duplicate column names if time_col accidentally appears in base_cols
    base_cols = [c for c in base_cols if c != time_col and c != target_col]

    df_feat = df_time[[time_col] + base_cols + [target_col]].copy()

    # Lag features
    for lag in lags:

```

```

    for c in base_cols:
        df_feat[f"{c}__lag{lag}"] = df_feat[c].shift(lag)

    # Rolling stats (past-only)
    for w in windows:
        for c in base_cols:
            df_feat[f"{c}__roll{w}_mean"] = (
                df_feat[c].shift(1).rolling(window=w, min_periods=w).mean()
            )
            df_feat[f"{c}__roll{w}_std"] = (
                df_feat[c].shift(1).rolling(window=w, min_periods=w).std()
            )

    # Deltas
    for c in base_cols:
        df_feat[f"{c}__delta1"] = df_feat[c].diff(1)

df_feat = df_feat.dropna().reset_index(drop=True)
return df_feat

base_feature_cols = [
    "Air temperature [K]",
    "Process temperature [K]",
    "Rotational speed [rpm]",
    "Torque [Nm]",
]

df_feat = make_time_features(
    df_tw,
    base_cols=base_feature_cols,
    time_col=TIME_COL,
    target_col=TARGET_COL,
    windows=(3, 5, 10),
    lags=(1, 2, 3),
)

print("Feature-engineered shape:", df_feat.shape)
print("Label distribution after feature engineering:")
print(df_feat[TARGET_COL].value_counts())
df_feat.head()

```

```

Feature-engineered shape: (236, 46)
Label distribution after feature engineering:
Machine failure
1      164
0       72
Name: count, dtype: int64

```

[5]: Tool wear [min] Air temperature [K] Process temperature [K] Rotational speed [rpm] Torque [Nm] \

0	11	300.000000	309.876190
1505.285714	41.078571		
1	12	300.188000	310.204000
1534.220000	41.368000		
2	13	299.670000	309.726000
1530.360000	40.040000		
3	14	300.180851	310.021277
1537.191489	38.055319		
4	15	299.658491	309.945283
1564.018868	39.549057		

Machine failure Air temperature [K]__lag1 Process temperature [K]__lag1
Rotational speed [rpm]__lag1 \

0	1	300.148889	310.191111
1515.666667			
1	1	300.000000	309.876190
1505.285714			
2	0	300.188000	310.204000
1534.220000			
3	0	299.670000	309.726000
1530.360000			
4	1	300.180851	310.021277
1537.191489			

Torque [Nm]__lag1 Air temperature [K]__lag2 Process temperature [K]__lag2
Rotational speed [rpm]__lag2 \

0	42.362222	300.196364	310.089091
1547.781818			
1	41.078571	300.148889	310.191111
1515.666667			
2	41.368000	300.000000	309.876190
1505.285714			
3	40.040000	300.188000	310.204000
1534.220000			
4	38.055319	299.670000	309.726000
1530.360000			

Torque [Nm]__lag2 Air temperature [K]__lag3 Process temperature [K]__lag3
Rotational speed [rpm]__lag3 \

0	39.861818	299.888889	309.827778
1513.111111			
1	42.362222	300.196364	310.089091
1547.781818			
2	41.078571	300.148889	310.191111
1515.666667			

3	41.368000	300.000000	309.876190
1505.285714			
4	40.040000	300.188000	310.204000
1534.220000			

	Torque [Nm] __lag3	Air temperature [K] __roll13_mean	Air temperature [K] __roll13_std
0	42.172222	300.078047	0.165527
1	39.861818	300.115084	0.102454
2	42.362222	300.112296	0.099198
3	41.078571	299.952667	0.262224
4	41.368000	300.012950	0.297025

	Process temperature [K] __roll13_mean	Process temperature [K] __roll13_std	Rotational speed [rpm] __roll13_mean
0	310.035993	0.187396	1525.519865
1	310.052131	0.160681	1522.911400
2	310.090434	0.185652	1518.390794
3	309.935397	0.244438	1523.288571
4	309.983759	0.241198	1533.923830

	Rotational speed [rpm] __roll13_std	Torque [Nm] __roll13_mean	Torque [Nm] __roll13_std
0	19.321714	41.465421	1.392006
1	22.155007	41.100871	1.250351
2	14.658236	41.602931	0.673301
3	15.709935	40.828857	0.698329
4	3.425361	39.821106	1.667153

	Air temperature [K] __roll15_mean	Air temperature [K] __roll15_std	Process temperature [K] __roll15_mean
0	300.014826	0.190990	

309.978036		
1	300.065794	0.122903
310.021662		
2	300.084428	0.134885
310.037634		
3	300.040651	0.221759
310.017278		
4	300.037548	0.219129
310.003716		

	Process temperature [K]__roll5_std	Rotational speed [rpm]__roll5_mean
	Rotational speed [rpm]__roll5_std \	
0	0.225854	1530.775323
17.056630		
1	0.160089	1522.051821
16.615627		
2	0.176057	1523.213062
17.357734		
3	0.209158	1526.662840
16.543037		
4	0.205504	1524.544774
13.574778		

	Torque [Nm]__roll5_mean	Torque [Nm]__roll5_std	Air temperature
	[K]__roll10_mean	Air temperature [K]__roll10_std \	
0	40.777195	1.558879	
299.990513		0.191651	
1	41.258070	1.028642	
299.994846		0.191291	
2	41.368567	0.998449	
299.986400		0.179150	
3	40.942122	1.024397	
299.985459		0.180958	
4	40.580823	1.636417	
300.003838		0.191305	

	Process temperature [K]__roll10_mean	Process temperature [K]__roll10_std
	Rotational speed [rpm]__roll10_mean \	
0	309.969939	0.172112
1538.922398		
1	309.961975	0.174661
1536.959303		
2	309.968172	0.182674
1534.829129		
3	309.958125	0.193760
1537.038659		
4	309.973194	0.192048

1538.169572

```
Rotational speed [rpm]__roll10_std  Torque [Nm]__roll10_mean  Torque
[Nm]__roll10_std  Air temperature [K]__delta1 \
0                  32.956080          40.302721
1.878917          -0.148889
1                  34.434572          40.344411
1.892343          0.188000
2                  33.811932          40.516574
1.900081          -0.518000
3                  32.582724          40.356162
1.861636          0.510851
4                  32.347893          40.049929
1.971109          -0.522360

Process temperature [K]__delta1  Rotational speed [rpm]__delta1  Torque
[Nm]__delta1
0                  -0.314921         -10.380952
-1.283651
1                  0.327810          28.934286
0.289429
2                  -0.478000         -3.860000
-1.328000
3                  0.295277          6.831489
-1.984681
4                  -0.075994         26.827379
1.493737
```

1.0.6 6. Train/test split (time-ordered)

We split by time order (no shuffle). If the default 80/20 split produces only one class in train or test, we automatically find the earliest valid split with positives in both sets.

```
[6]: X = df_feat.drop(columns=[TARGET_COL])
y = df_feat[TARGET_COL].astype(int)

print("Overall label distribution:")
print(y.value_counts())

def find_valid_split(y_series: pd.Series, min_pos_train=1, min_pos_test=1, ↴
                     min_train_size=50):
    y_series = y_series.reset_index(drop=True)
    total_pos = int(y_series.sum())
    if total_pos < (min_pos_train + min_pos_test):
        raise ValueError(
            f"Not enough positive samples overall (found {total_pos}) "
            f"to have positives in both train and test."
```

```

        )
for split_idx in range(min_train_size, len(y_series) - 1):
    pos_train = int(y_series.iloc[:split_idx].sum())
    pos_test = int(y_series.iloc[split_idx:].sum())
    if pos_train >= min_pos_train and pos_test >= min_pos_test:
        return split_idx
raise ValueError("Could not find a time-ordered split with positives in both sets.")

default_split = int(len(df_feat) * 0.8)
if y.iloc[:default_split].nunique() < 2 or y.iloc[default_split:].nunique() < 2:
    split_idx = find_valid_split(y, min_pos_train=1, min_pos_test=1, min_train_size=50)
else:
    split_idx = default_split

X_train, X_test = X.iloc[:split_idx], X.iloc[split_idx:]
y_train, y_test = y.iloc[:split_idx], y.iloc[split_idx:]

print("Split index:", split_idx)
print("Train label counts:\n", y_train.value_counts())
print("Test label counts:\n", y_test.value_counts())

```

Overall label distribution:

Machine failure

1	164
0	72

Name: count, dtype: int64

Split index: 188

Train label counts:

Machine failure

1	126
0	62

Name: count, dtype: int64

Test label counts:

Machine failure

1	38
0	10

Name: count, dtype: int64

1.0.7 7. Baseline model: Logistic Regression

Standardized Logistic Regression with `class_weight='balanced'`.

```
[7]: logreg = Pipeline(steps=[
    ("scaler", StandardScaler()),
    ("clf", LogisticRegression(max_iter=2000, class_weight="balanced"))
])
```

```

logreg.fit(X_train, y_train)

proba_test_lr = logreg.predict_proba(X_test)[:, 1]
pred_test_lr = (proba_test_lr >= 0.5).astype(int)

print("Logistic Regression")
print("ROC-AUC:", round(roc_auc_score(y_test, proba_test_lr), 4))
print("PR-AUC :", round(average_precision_score(y_test, proba_test_lr), 4))
print()
print(classification_report(y_test, pred_test_lr, digits=4))
print("Confusion matrix:\n", confusion_matrix(y_test, pred_test_lr))

```

Logistic Regression

ROC-AUC: 0.5947

PR-AUC : 0.8307

	precision	recall	f1-score	support
0	0.1875	0.6000	0.2857	10
1	0.7500	0.3158	0.4444	38
accuracy			0.3750	48
macro avg	0.4688	0.4579	0.3651	48
weighted avg	0.6328	0.3750	0.4114	48

Confusion matrix:

```

[[ 6  4]
[26 12]]

```

1.0.8 8. Tree-based model: Random Forest and Gradient Boosting

Tree-based models can capture non-linear interactions without requiring feature scaling.

```

[8]: rf = RandomForestClassifier(
    n_estimators=400,
    min_samples_split=10,
    min_samples_leaf=5,
    class_weight="balanced",
    random_state=42,
    n_jobs=-1
)

gb = GradientBoostingClassifier(random_state=42)

rf.fit(X_train, y_train)
gb.fit(X_train, y_train)

```

```

proba_test_rf = rf.predict_proba(X_test)[:, 1]
pred_test_rf = (proba_test_rf >= 0.5).astype(int)

proba_test_gb = gb.predict_proba(X_test)[:, 1]
pred_test_gb = (proba_test_gb >= 0.5).astype(int)

def print_scores(name, proba, pred):
    print(name)
    print("ROC-AUC:", round(roc_auc_score(y_test, proba), 4))
    print("PR-AUC :", round(average_precision_score(y_test, proba), 4))
    print(classification_report(y_test, pred, digits=4))
    print("Confusion matrix:\n", confusion_matrix(y_test, pred))
    print("-"*70)

print_scores("Random Forest", proba_test_rf, pred_test_rf)
print_scores("Gradient Boosting", proba_test_gb, pred_test_gb)

```

Random Forest

ROC-AUC: 0.5237

PR-AUC : 0.8225

	precision	recall	f1-score	support
0	0.2308	0.3000	0.2609	10
1	0.8000	0.7368	0.7671	38
accuracy			0.6458	48
macro avg	0.5154	0.5184	0.5140	48
weighted avg	0.6814	0.6458	0.6617	48

Confusion matrix:

```
[[ 3  7]
 [10 28]]
```

Gradient Boosting

ROC-AUC: 0.7342

PR-AUC : 0.8993

	precision	recall	f1-score	support
0	0.2432	0.9000	0.3830	10
1	0.9091	0.2632	0.4082	38
accuracy			0.3958	48
macro avg	0.5762	0.5816	0.3956	48
weighted avg	0.7704	0.3958	0.4029	48

Confusion matrix:

```
[[ 9  1]
 [28 10]]
```

1.0.9 9. Evaluation plots

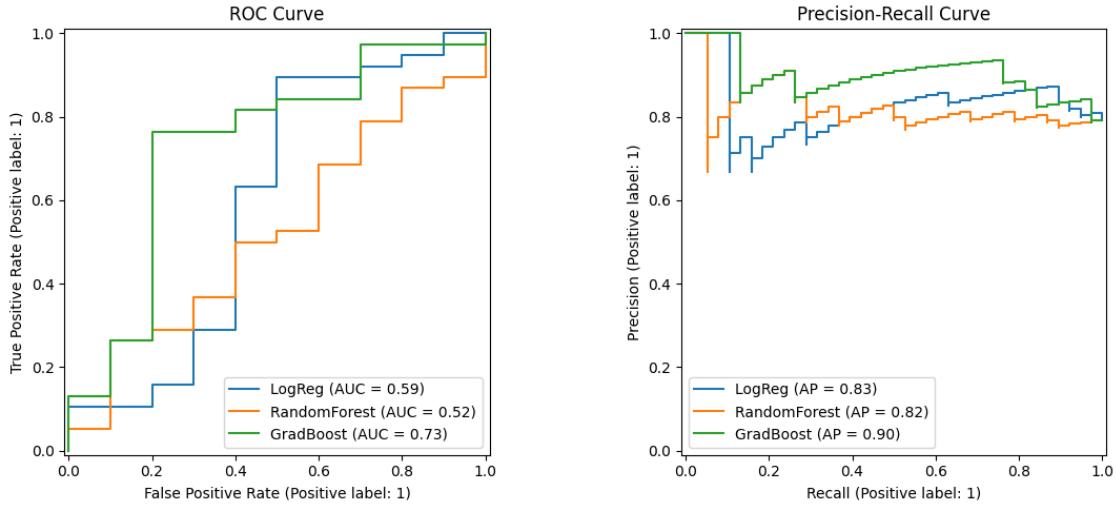
ROC and Precision-Recall curves help compare performance under class imbalance.

```
[9]: fig, axes = plt.subplots(1, 2, figsize=(12, 5))

# ROC Curve (left)
RocCurveDisplay.from_predictions(
    y_test, proba_test_lr, name="LogReg", ax=axes[0]
)
RocCurveDisplay.from_predictions(
    y_test, proba_test_rf, name="RandomForest", ax=axes[0]
)
RocCurveDisplay.from_predictions(
    y_test, proba_test_gb, name="GradBoost", ax=axes[0]
)
axes[0].set_title("ROC Curve")

# Precision-Recall Curve (right)
PrecisionRecallDisplay.from_predictions(
    y_test, proba_test_lr, name="LogReg", ax=axes[1]
)
PrecisionRecallDisplay.from_predictions(
    y_test, proba_test_rf, name="RandomForest", ax=axes[1]
)
PrecisionRecallDisplay.from_predictions(
    y_test, proba_test_gb, name="GradBoost", ax=axes[1]
)
axes[1].set_title("Precision-Recall Curve")

plt.tight_layout()
plt.show()
```



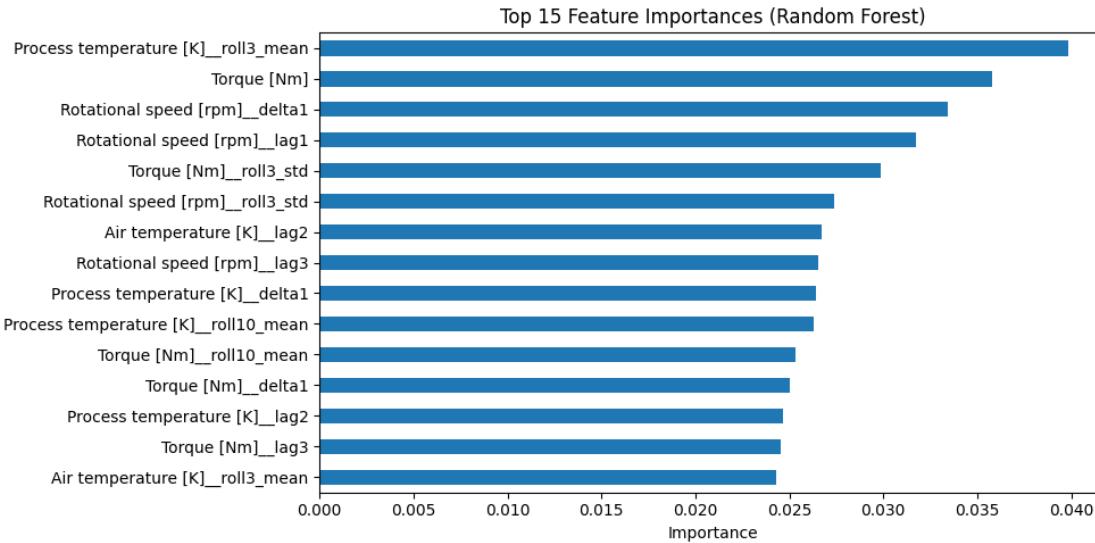
1.0.10 10. Feature importance (tree-based)

Random Forest feature importance provides a quick view of what the model uses most.

```
[10]: importances = pd.Series(rf.feature_importances_, index=X_train.columns).
       ↪sort_values(ascending=False)
top_k = 15
top = importances.head(top_k)

plt.figure(figsize=(10,5))
top[::-1].plot(kind="barh")
plt.title(f"Top {top_k} Feature Importances (Random Forest)")
plt.xlabel("Importance")
plt.tight_layout()
plt.show()

top.to_frame("importance")
```



[10] :

	importance
Process temperature [K]_roll3_mean	0.039847
Torque [Nm]	0.035810
Rotational speed [rpm]_delta1	0.033444
Rotational speed [rpm]_lag1	0.031755
Torque [Nm]_roll3_std	0.029845
Rotational speed [rpm]_roll3_std	0.027405
Air temperature [K]_lag2	0.026741
Rotational speed [rpm]_lag3	0.026511
Process temperature [K]_delta1	0.026407
Process temperature [K]_roll10_mean	0.026291
Torque [Nm]_roll10_mean	0.025317
Torque [Nm]_delta1	0.025014
Process temperature [K]_lag2	0.024666
Torque [Nm]_lag3	0.024558
Air temperature [K]_roll3_mean	0.024290

1.0.11 11. Save model-ready artifacts

This saves the engineered time-series dataset used by Model 1 for reproducibility.

```
[11]: OUTPUT_FEATURES_PATH = "../data/ai4i_time_features_model1.csv"
df_feat.to_csv(OUTPUT_FEATURES_PATH, index=False)
print(f"Saved time-feature dataset to: {OUTPUT_FEATURES_PATH}")
```

Saved time-feature dataset to: ../data/ai4i_time_features_model1.csv

1.0.12 12. Export Classification Results for Tableau

```
[15]: os.makedirs("../outputs", exist_ok=True)

# Collect available probability vectors from this notebook
candidates = []
if "proba_test_lr" in globals():
    candidates["LogReg"] = np.asarray(proba_test_lr)
if "proba_test_rf" in globals():
    candidates["RandomForest"] = np.asarray(proba_test_rf)
if "proba_test_gb" in globals():
    candidates["GradBoost"] = np.asarray(proba_test_gb)

if not candidates:
    raise NameError("No probability vectors found. Expected one of: " +
                   "proba_test_lr, proba_test_rf, proba_test_gb.")

# Pick best model by ROC-AUC on test set
scores = {name: roc_auc_score(y_test, proba) for name, proba in candidates.items()}
best_name = max(scores, key=scores.get)
y_pred_prob = candidates[best_name]

print("ROC-AUC by model:", {k: round(v, 4) for k, v in scores.items()})
print(f"Selected for export: {best_name}")

# Recommended threshold (Youden's J)
fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob)
optimal_idx = (tpr - fpr).argmax()
recommended_threshold = float(thresholds[optimal_idx])
print(f"Recommended threshold (Youden's J): {recommended_threshold:.6f}")

# Pred label using recommended threshold
y_pred_label = (y_pred_prob >= recommended_threshold).astype(int)

# Use notebook's TIME_COL if available
time_col = TIME_COL if "TIME_COL" in globals() else "Tool wear [min]"
if time_col not in X_test.columns:
    raise KeyError(f'{time_col}' not found in X_test columns: {list(X_test.columns)})"

export_df = pd.DataFrame({
    "tool_wear": X_test[time_col].values,
    "actual_failure": np.asarray(y_test).astype(int),
    "pred_prob_failure": y_pred_prob,
    "pred_label": y_pred_label,
    "model_selected": best_name,
```

```

    "recommended_threshold": recommended_threshold
}).sort_values("tool_wear")

export_path = "../outputs/pred_failure_traditional.csv"
export_df.to_csv(export_path, index=False)

print(f"Saved: {export_path} | rows={len(export_df)}")

```

ROC-AUC by model: {'LogReg': 0.5947, 'RandomForest': 0.5237, 'GradBoost': 0.7342}
Selected for export: GradBoost
Recommended threshold (Youden's J): 0.016185
Saved: ../outputs/pred_failure_traditional.csv | rows=48

1.0.13 Why We Chose Failure Classification Instead of Torque Prediction

In an earlier draft, we considered predicting future torque values as a regression task. However, after reviewing the project objectives, we shifted to predicting **Machine failure (binary classification)** for the following reasons:

1. Alignment with Project Goal

The primary objective of predictive maintenance is to anticipate equipment failure, not merely forecast sensor values. Predicting failure directly better reflects real-world maintenance decision-making.

2. Business Relevance

Maintenance teams act on failure risk (fail vs. no-fail), not on small changes in torque. A classification model provides actionable outputs such as failure probability and risk thresholds.

3. Clear Evaluation Metrics

Failure classification allows the use of metrics suited for rare-event detection (ROC-AUC, Precision-Recall AUC), which are more appropriate for imbalanced industrial datasets.

4. Avoiding Indirect Modeling

Predicting torque and then inferring failure from torque changes adds an extra modeling step and potential error propagation. Directly modeling failure simplifies interpretation and reduces complexity.

For these reasons, we adopted a time-aware classification framework to predict machine failure using engineered temporal features derived from sensor data.

model1B_regression_torque_forecast

February 21, 2026

1 Model 1B — Regression Baseline (Torque Forecasting)

This notebook is an **optional baseline** included to contrast regression vs. classification for the project.

Goal (regression): - Predict the **next-step torque** value (`Torque_future`) using an ordered sequence derived from `Tool_wear [min]`.

(*Predict mean torque at `tool_wear = t+1` using features at `tool_wear = t`*)

Important note: - This regression task is *not* the primary predictive maintenance objective (failure prediction). - It is included to demonstrate why **direct failure classification** is more actionable for maintenance decisions.

1.0.1 1. Setup

```
[7]: import os
import warnings
warnings.filterwarnings("ignore")

import numpy as np
import pandas as pd

from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression, Ridge
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

import matplotlib.pyplot as plt

pd.set_option("display.max_columns", 200)
pd.set_option("display.width", 120)
```

1.0.2 2. Load data

We load the prepared dataset if available. Otherwise, we load the raw dataset and apply minimal preprocessing (drop identifiers and one-hot encode Type).

For regression, the failure-mode flags are not used.

```
[8]: DATA_PREPARED_PATH = "../data/ai4i_prepared.csv"
DATA_RAW_PATH = "../data/ai4i_2020_predictive_maintenance.csv"

TIME_COL = "Tool wear [min]"

def load_dataset():
    if os.path.exists(DATA_PREPARED_PATH):
        df = pd.read_csv(DATA_PREPARED_PATH)
        source = "prepared"
    else:
        df_raw = pd.read_csv(DATA_RAW_PATH)

        # Minimal preprocessing
        drop_cols = ["UDI", "Product ID"]
        df = df_raw.drop(columns=[c for c in drop_cols if c in df_raw.columns], errors="ignore")

        if "Type" in df.columns:
            df = pd.get_dummies(df, columns=["Type"], drop_first=True)

        source = "raw+prepped"
    return df, source

df, source = load_dataset()
print(f"Loaded source: {source}")
print("Shape:", df.shape)
df.head()
```

Loaded source: prepared

Shape: (10000, 14)

	UDI	Product ID	Type	Air temperature [K]	Process temperature [K]
				Rotational speed [rpm]	Torque [Nm]
0	1	M14860	M		298.1
1551				42.8	
1	7257	H36670	H		300.2
1408				42.5	
2	504	M15363	M		297.6
1442				48.1	
3	7169	L54348	L		300.3
1704				29.5	
4	7089	M21948	M		300.6
1614				32.7	
				Tool wear [min]	Machine failure
0				0	TWF
1				0	HDF
				0	PWF
				0	OSF
				0	RNF

2	0	0	0	0	0	0
3	0	0	0	0	0	0
4	0	0	0	0	0	0

1.0.3 3. Create an ordered sequence using tool wear (time proxy)

We aggregate sensor readings by Tool wear [min] to form a monotonic sequence. Then we define Torque_future as the next step torque value (one-step-ahead forecasting).

```
[9]: required_cols = [
    "Air temperature [K]",
    "Process temperature [K]",
    "Rotational speed [rpm]",
    "Torque [Nm]",
    TIME_COL,
]
missing = [c for c in required_cols if c not in df.columns]
if missing:
    raise ValueError(f"Missing required columns: {missing}")

# Aggregate by tool wear (population-level) to build an ordered sequence
df_tw = (
    df
    .groupby(TIME_COL, as_index=False)
    .mean(numeric_only=True)
    .sort_values(TIME_COL)
    .reset_index(drop=True)
)

# One-step-ahead target
df_tw["Torque_future"] = df_tw["Torque [Nm]"].shift(-1)
df_tw = df_tw.dropna().reset_index(drop=True)

print("Aggregated shape:", df_tw.shape)
df_tw.head()
```

Aggregated shape: (245, 13)

```
[9]:   Tool wear [min]          UDI Air temperature [K]  Process temperature [K]
  Rotational speed [rpm]  Torque [Nm] \
0           0 5014.266667          299.956667          309.955833
1524.916667  40.661667
1           2 5038.826087          300.272464          310.142029
1555.521739  39.646377
2           3 4925.970588          299.679412          309.826471
1508.264706  41.644118
3           4 5505.205882          299.997059          309.870588
1525.882353  41.117647
```

4	5	4937.206349	299.925397	310.014286
1620.761905	36.071429			

	Machine failure	TWF	HDF	PWF	OSF	RNF	Torque_future
0	0.025000	0.0	0.000000	0.025000	0.0	0.000000	39.646377
1	0.028986	0.0	0.014493	0.014493	0.0	0.014493	41.644118
2	0.029412	0.0	0.000000	0.029412	0.0	0.000000	41.117647
3	0.000000	0.0	0.000000	0.000000	0.0	0.000000	36.071429
4	0.015873	0.0	0.000000	0.015873	0.0	0.000000	38.674194

1.0.4 4. Train/test split (time-ordered)

We split without shuffling to preserve time order.

```
[10]: feature_cols = [
    "Air temperature [K]",
    "Process temperature [K]",
    "Rotational speed [rpm]",
    TIME_COL,
    "Torque [Nm]", # current torque can help predict next torque
]

X = df_tw[feature_cols]
y = df_tw["Torque_future"]

split_idx = int(len(df_tw) * 0.8)
X_train, X_test = X.iloc[:split_idx], X.iloc[split_idx:]
y_train, y_test = y.iloc[:split_idx], y.iloc[split_idx:]

print("Train size:", X_train.shape, " Test size:", X_test.shape)
```

Train size: (196, 5) Test size: (49, 5)

1.0.5 5. Regression models

We compare a simple linear baseline vs. a non-linear model:

- Linear Regression (with scaling)
- Ridge Regression (regularized linear)
- Random Forest Regressor (non-linear baseline)

```
[11]: models = {}

models["LinearRegression"] = Pipeline(steps=[
    ("scaler", StandardScaler()),
    ("reg", LinearRegression())
])

models["Ridge"] = Pipeline(steps=[
    ("scaler", StandardScaler()),
    ("reg", Ridge(alpha=1.0, random_state=42))
```

```

])

models["RandomForestRegressor"] = RandomForestRegressor(
    n_estimators=400,
    min_samples_split=10,
    min_samples_leaf=5,
    random_state=42,
    n_jobs=-1
)

def eval_regression(y_true, y_pred):
    mae = mean_absolute_error(y_true, y_pred)
    mse = mean_squared_error(y_true, y_pred)
    rmse = np.sqrt(mse)
    r2 = r2_score(y_true, y_pred)
    return mae, rmse, r2

results = []
preds = {}

for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    preds[name] = y_pred
    mae, rmse, r2 = eval_regression(y_test, y_pred)
    results.append((name, mae, rmse, r2))

results_df = pd.DataFrame(results, columns=["model", "MAE", "RMSE", "R2"]).
    ↪sort_values("RMSE")
results_df

```

```
[11]:      model      MAE      RMSE      R2
1          Ridge  3.145049  4.524447 -0.021235
0  LinearRegression  3.143357  4.524576 -0.021293
2 RandomForestRegressor  3.177763  4.525682 -0.021792
```

1.0.6 6. Plot: actual vs predicted torque (test set)

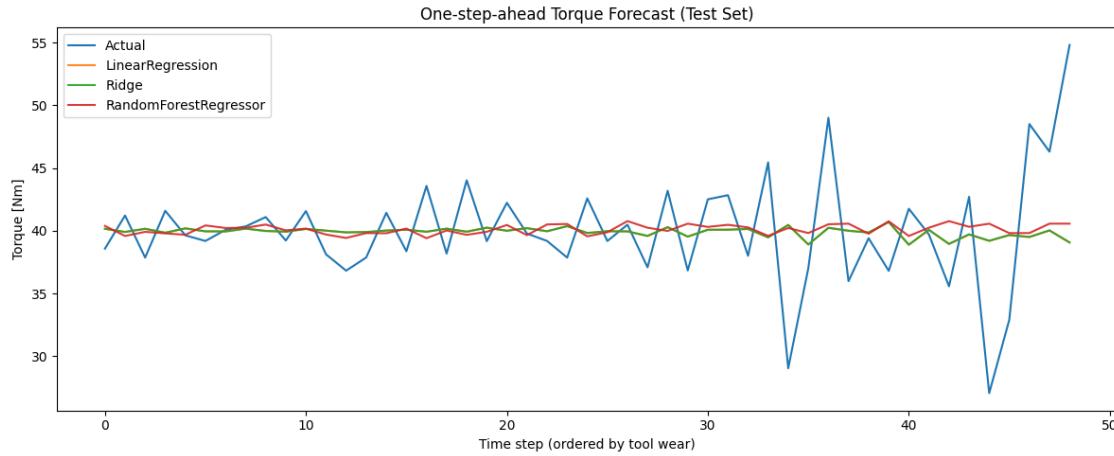
This visual comparison helps show how well each model tracks the next-step torque values.

```
[12]: plt.figure(figsize=(12, 5))
plt.plot(y_test.reset_index(drop=True).values, label="Actual")
for name, y_pred in preds.items():
    plt.plot(y_pred, label=name, alpha=0.9)
plt.title("One-step-ahead Torque Forecast (Test Set)")
plt.xlabel("Time step (ordered by tool wear)")
plt.ylabel("Torque [Nm]")
```

```

plt.legend()
plt.tight_layout()
plt.show()

```



1.0.7 7. Discussion: regression vs. failure classification

- Torque forecasting predicts a continuous sensor value, which may be useful for anomaly detection.
- Predictive maintenance decisions typically require **failure risk** (classification) rather than a raw sensor forecast.
- Even if torque is predicted accurately, mapping that forecast to a failure decision adds an extra step and can amplify errors.

1.0.8 8. Save predictions

This exports a small table with actual and predicted values for reporting.

```

[13]: out = pd.DataFrame({
    "Torque_future_actual": y_test.values,
    **{f"Torque_future_pred_{name}": y_pred for name, y_pred in preds.items()}
})

OUTPUT_PATH = "../data/model1B_torque_forecast_predictions.csv"
out.to_csv(OUTPUT_PATH, index=False)
print(f"Saved predictions to: {OUTPUT_PATH}")
out.head()

```

Saved predictions to: ../data/model1B_torque_forecast_predictions.csv

	Torque_future_actual	Torque_future_pred_LinearRegression
Torque_future_pred_Ridge	\	
0	38.569565	40.150744

```

40.149124
1           41.208511          39.899440
39.898446
2           37.856410          40.151312
40.150230
3           41.584444          39.837399
39.837839
4           39.628889          40.185983
40.183500

```

```

    Torque_future_pred_RandomForestRegressor
0                  40.381197
1                  39.578778
2                  39.907429
3                  39.786314
4                  39.674121

```

1.0.9 9. Export Regression Forecast Results for Tableau

```
[16]: os.makedirs("../outputs", exist_ok=True)

# Validate preds exists
if "preds" not in globals() or not isinstance(preds, dict) or len(preds) == 0:
    raise NameError("preds dict not found (or empty). Run the model training/
    ↪evaluation cell first.")

# Pick best model (lowest RMSE)
if "results_df" in globals() and hasattr(results_df, "iloc") and "model" in_
↪results_df.columns:
    best_name = results_df.iloc[0]["model"]
else:
    rmse_scores = {name: float(np.sqrt(mean_squared_error(y_test, pred))) for_
    ↪name, pred in preds.items()}
    best_name = min(rmse_scores, key=rmse_scores.get)
    print("RMSE by model:", {k: round(v, 4) for k, v in rmse_scores.items()})

if best_name not in preds:
    raise KeyError(f"Best model '{best_name}' not found in preds keys:_
    ↪{list(preds.keys())}")

y_pred = np.asarray(preds[best_name]).ravel()

print(f"Selected for export: {best_name}")

# Use the notebook's time column
time_col = TIME_COL if "TIME_COL" in globals() else "Tool wear [min]"
if time_col not in X_test.columns:
```

```

    raise KeyError(f"'{time_col}' not found in X_test columns: {list(X_test.
    ↪columns)}")

# Build export dataframe
export_df = pd.DataFrame({
    "tool_wear": X_test[time_col].values,
    "torque_actual_next": np.asarray(y_test).ravel(),
    "torque_pred": y_pred,
    "model_selected": best_name
}).sort_values("tool_wear")

export_df["residual"] = export_df["torque_actual_next"] - □
    ↪export_df["torque_pred"]

# Save
export_path = "../outputs/pred_torque_forecast.csv"
export_df.to_csv(export_path, index=False)

print(f"Saved: {export_path} | rows={len(export_df)}")

```

Selected for export: Ridge
 Saved: ../outputs/pred_torque_forecast.csv | rows=49

model2_deep_learning

February 21, 2026

1 Model 2 — Deep Learning (Sequence Model using Tool-Wear Time Proxy)

This notebook implements **Model 2** for the project using deep learning.

Project framing: - Dataset: AI4I 2020 Predictive Maintenance - Primary target: `Machine failure` (binary classification) - Time proxy: `Tool wear [min]` is treated as progression over time-under-use.

Key design choices: 1. We aggregate records by tool wear to create an ordered sequence (population-level). 2. We preserve the binary label during aggregation using `max()` for `Machine failure`. 3. We build sliding-window sequences and train a sequence model (LSTM).

Deliverables in this notebook: - Data preparation for deep learning (sequence windows) - Baseline deep model (MLP) and sequence model (LSTM) - Time-aware train/validation/test split (no shuffle) - Evaluation with ROC-AUC and PR-AUC (useful for class imbalance)

1.0.1 1. Setup

```
[1]: import os
import warnings
warnings.filterwarnings("ignore")

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.preprocessing import StandardScaler
from sklearn.utils.class_weight import compute_class_weight
from sklearn.metrics import (
    classification_report,
    confusion_matrix,
    roc_auc_score,
    roc_curve,
    average_precision_score,
    RocCurveDisplay,
    PrecisionRecallDisplay
)
```

```

pd.set_option("display.max_columns", 200)
pd.set_option("display.width", 120)

# Deep learning (TensorFlow / Keras)
try:
    import tensorflow as tf
    from tensorflow import keras
    from tensorflow.keras import layers
    print("TensorFlow:", tf.__version__)
except ModuleNotFoundError as e:
    raise ModuleNotFoundError(
        "TensorFlow is not installed. Add `tensorflow>=2.12` to requirements."
        "txt and reinstall."
    ) from e

```

TensorFlow: 2.16.2

1.0.2 2. Load data

We load the prepared dataset if available. Otherwise, we load raw CSV and apply minimal preprocessing:

- Drop identifiers (UDI, Product ID)
- Drop failure-mode flags (TWF, HDF, PWF, OSF, RNF) to avoid leakage
- One-hot encode Type

```

[2]: DATA_PREPARED_PATH = "../data/ai4i_prepared.csv"
DATA_RAW_PATH = "../data/ai4i_2020_predictive_maintenance.csv"

TARGET_COL = "Machine failure"
TIME_COL = "Tool wear [min]"

def load_dataset():
    if os.path.exists(DATA_PREPARED_PATH):
        df = pd.read_csv(DATA_PREPARED_PATH)
        source = "prepared"
    else:
        df_raw = pd.read_csv(DATA_RAW_PATH)
        drop_cols = ["UDI", "Product ID", "TWF", "HDF", "PWF", "OSF", "RNF"]
        df = df_raw.drop(columns=[c for c in drop_cols if c in df_raw.columns], u
        ↪errors="ignore")
        if "Type" in df.columns:
            df = pd.get_dummies(df, columns=["Type"], drop_first=True)
        source = "raw+prepped"
    return df, source

df, source = load_dataset()
print(f"Loaded source: {source}")
print("Shape:", df.shape)
df.head()

```

```
Loaded source: prepared  
Shape: (10000, 14)
```

```
[2]:    UDI Product ID Type Air temperature [K] Process temperature [K]  
Rotational speed [rpm] Torque [Nm] \\\n0     1      M14860   M                 298.1                  308.6  
1551          42.8  
1    7257      H36670   H                 300.2                  310.3  
1408          42.5  
2    504       M15363   M                 297.6                  309.2  
1442          48.1  
3    7169      L54348   L                 300.3                  310.3  
1704          29.5  
4    7089      M21948   M                 300.6                  310.3  
1614          32.7
```


Tool wear [min]	Machine failure	TWF	HDF	PWF	OSF	RNF
0	0	0	0	0	0	0
1	0	0	0	0	0	0
2	0	0	0	0	0	0
3	0	0	0	0	0	0
4	0	0	0	0	0	0

1.0.3 3. Aggregate by tool wear to form an ordered sequence

Deep learning sequence models need ordered samples. We aggregate by tool wear.

Important: Machine failure is aggregated with `max()` to preserve binary semantics.

```
[3]: required_cols = [  
    "Air temperature [K]",  
    "Process temperature [K]",  
    "Rotational speed [rpm]",  
    "Torque [Nm]",  
    TIME_COL,  
    TARGET_COL  
]  
missing = [c for c in required_cols if c not in df.columns]  
if missing:  
    raise ValueError(f"Missing required columns: {missing}")  
  
df_tw = (  
    df  
    .groupby(TIME_COL, as_index=False)  
    .agg({  
        "Air temperature [K]": "mean",  
        "Process temperature [K]": "mean",  
        "Rotational speed [rpm]": "mean",  
        "Machine failure": "max"}))
```

```

        "Torque [Nm]": "mean",
        TARGET_COL: "max",
    })
    .sort_values(TIME_COL)
    .reset_index(drop=True)
)

print("Aggregated shape:", df_tw.shape)
print("Label distribution:", df_tw[TARGET_COL].value_counts().to_dict())
df_tw.head()

```

Aggregated shape: (246, 6)
Label distribution: {1: 172, 0: 74}

[3]:

	Tool wear [min]	Air temperature [K]	Process temperature [K]	Rotational speed [rpm]	Torque [Nm]	Machine failure
0		299.956667		309.955833	0	
1524.916667	40.661667		1		2	300.272464
1		300.272464		310.142029	1	
1555.521739	39.646377		1		3	299.679412
2		299.679412		309.826471	1	
1508.264706	41.644118		1		4	299.997059
3		299.997059		309.870588	0	
1525.882353	41.117647		0		5	299.925397
4		299.925397		310.014286	1	
1620.761905	36.071429					

1.0.4 4. Build sliding-window sequences

We create sequences of length SEQ_LEN from the ordered tool-wear series. Each sequence uses sensor readings from the previous SEQ_LEN steps. The label for a sequence is the `Machine failure` value at the final step of the window.

[4]:

```

# Features used for the sequence model
feature_cols = [
    "Air temperature [K]",
    "Process temperature [K]",
    "Rotational speed [rpm]",
    "Torque [Nm]",
]

SEQ_LEN = 10 # window length (can be tuned)

X_values = df_tw[feature_cols].values.astype(np.float32)
y_values = df_tw[TARGET_COL].values.astype(int)

def make_sequences(X, y, seq_len: int):
    X_seq, y_seq = [], []

```

```

    for i in range(seq_len - 1, len(X)):
        X_seq.append(X[i - seq_len + 1 : i + 1])
        y_seq.append(y[i]) # label at window end
    return np.array(X_seq, dtype=np.float32), np.array(y_seq, dtype=int)

X_seq, y_seq = make_sequences(X_values, y_values, SEQ_LEN)

print("X_seq shape:", X_seq.shape, "(samples, timesteps, features)")
print("y_seq distribution:", dict(pd.Series(y_seq).value_counts()))

```

X_seq shape: (237, 10, 4) (samples, timesteps, features)
y_seq distribution: {1: 165, 0: 72}

1.0.5 5. Time-aware train/validation/test split

We split sequences by time order (no shuffle). We also ensure that the training set contains both classes. Because failures can be rare, we select the earliest split points that keep positives in all sets.

```
[5]: def find_splits_with_positives(y, min_pos_train=1, min_pos_val=1, min_pos_test=1, min_train=50, min_val=20):
    n = len(y)
    y = pd.Series(y).reset_index(drop=True)

    total_pos = int(y.sum())
    if total_pos < (min_pos_train + min_pos_val + min_pos_test):
        raise ValueError(f"Not enough positive samples overall (found {total_pos}).")

    # Search for train_end and val_end (time-ordered)
    for train_end in range(min_train, n - (min_val + 1)):
        pos_train = int(y.iloc[:train_end].sum())
        if pos_train < min_pos_train:
            continue
        for val_end in range(train_end + min_val, n - 1):
            pos_val = int(y.iloc[train_end:val_end].sum())
            pos_test = int(y.iloc[val_end:].sum())
            if pos_val >= min_pos_val and pos_test >= min_pos_test:
                return train_end, val_end

    raise ValueError("Could not find time-ordered splits with positives in train/val/test.")

train_end, val_end = find_splits_with_positives(y_seq)

X_train, y_train = X_seq[:train_end], y_seq[:train_end]
X_val, y_val = X_seq[train_end:val_end], y_seq[train_end:val_end]
X_test, y_test = X_seq[val_end:], y_seq[val_end:]
```

```

print("Split indices:", train_end, val_end)
print("Train:", X_train.shape, "labels:", dict(pd.Series(y_train).
    value_counts()))
print("Val  :", X_val.shape,   "labels:", dict(pd.Series(y_val).value_counts()))
print("Test :", X_test.shape,  "labels:", dict(pd.Series(y_test).
    value_counts()))

```

```

Split indices: 50 70
Train: (50, 10, 4) labels: {1: 32, 0: 18}
Val  : (20, 10, 4) labels: {1: 16, 0: 4}
Test : (167, 10, 4) labels: {1: 117, 0: 50}

```

1.0.6 6. Feature scaling (fit on train only)

We standardize features using statistics from the training set only, then apply to validation and test. For sequences, we fit the scaler on the flattened training data and reshape back.

```

[6]: scaler = StandardScaler()

# Fit scaler on training data (flatten timesteps)
X_train_flat = X_train.reshape(-1, X_train.shape[-1])
scaler.fit(X_train_flat)

def scale_sequences(X, scaler_obj):
    X_flat = X.reshape(-1, X.shape[-1])
    X_scaled = scaler_obj.transform(X_flat)
    return X_scaled.reshape(X.shape)

X_train_s = scale_sequences(X_train, scaler)
X_val_s  = scale_sequences(X_val, scaler)
X_test_s = scale_sequences(X_test, scaler)

print("Scaled shapes:", X_train_s.shape, X_val_s.shape, X_test_s.shape)

```

```

Scaled shapes: (50, 10, 4) (20, 10, 4) (167, 10, 4)

```

1.0.7 7. Handle class imbalance (class weights)

We compute class weights from the training labels and pass them to Keras during training.

```

[7]: classes = np.unique(y_train)
class_weights = compute_class_weight(class_weight="balanced", classes=classes,
    y=y_train)
class_weight_dict = {int(c): float(w) for c, w in zip(classes, class_weights)}

print("Class weights:", class_weight_dict)

```

```

Class weights: {0: 1.3888888888888888, 1: 0.78125}

```

1.0.8 8. Model A (baseline deep model): MLP on last timestep

This baseline uses only the last timestep features from each sequence and trains a small feedforward network. It provides a deep-learning baseline that is comparable to traditional ML on tabular features.

```
[8]: # Use last timestep only (tabular baseline)
X_train_last = X_train_s[:, -1, :]
X_val_last   = X_val_s[:, -1, :]
X_test_last  = X_test_s[:, -1, :]

def build_mlp(input_dim: int):
    model = keras.Sequential([
        layers.Input(shape=(input_dim,)),
        layers.Dense(32, activation="relu"),
        layers.Dropout(0.2),
        layers.Dense(16, activation="relu"),
        layers.Dropout(0.2),
        layers.Dense(1, activation="sigmoid"),
    ])
    model.compile(
        optimizer=keras.optimizers.Adam(learning_rate=1e-3),
        loss="binary_crossentropy",
        metrics=[keras.metrics.AUC(name="roc_auc"), keras.metrics.
        ↪AUC(curve="PR", name="pr_auc")]
    )
    return model

mlp = build_mlp(X_train_last.shape[1])
mlp.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 32)	160
dropout (Dropout)	(None, 32)	0
dense_1 (Dense)	(None, 16)	528
dropout_1 (Dropout)	(None, 16)	0
dense_2 (Dense)	(None, 1)	17

```
Total params: 705 (2.75 KB)
```

```
Trainable params: 705 (2.75 KB)
```

```
Non-trainable params: 0 (0.00 B)
```

```
[9]: callbacks = [
    keras.callbacks.EarlyStopping(monitor="val_pr_auc", mode="max", patience=10, restore_best_weights=True),
    keras.callbacks.ReduceLROnPlateau(monitor="val_pr_auc", mode="max", factor=0.5, patience=5, min_lr=1e-5),
]

history_mlp = mlp.fit(
    X_train_last, y_train,
    validation_data=(X_val_last, y_val),
    epochs=80,
    batch_size=32,
    class_weight=class_weight_dict,
    callbacks=callbacks,
    verbose=1
)
```

```
Epoch 1/80
2/2          1s 153ms/step - loss:
0.7998 - pr_auc: 0.6314 - roc_auc: 0.4722 - val_loss: 1.0554 - val_pr_auc:
0.6516 - val_roc_auc: 0.1562 - learning_rate: 0.0010
Epoch 2/80
2/2          0s 24ms/step - loss:
0.7944 - pr_auc: 0.7019 - roc_auc: 0.4983 - val_loss: 1.0309 - val_pr_auc:
0.6609 - val_roc_auc: 0.1797 - learning_rate: 0.0010
Epoch 3/80
2/2          0s 25ms/step - loss:
0.7689 - pr_auc: 0.6679 - roc_auc: 0.5208 - val_loss: 1.0069 - val_pr_auc:
0.6634 - val_roc_auc: 0.1875 - learning_rate: 0.0010
Epoch 4/80
2/2          0s 26ms/step - loss:
0.7731 - pr_auc: 0.6123 - roc_auc: 0.4271 - val_loss: 0.9849 - val_pr_auc:
0.6664 - val_roc_auc: 0.1953 - learning_rate: 0.0010
Epoch 5/80
2/2          0s 25ms/step - loss:
0.7410 - pr_auc: 0.7070 - roc_auc: 0.5382 - val_loss: 0.9656 - val_pr_auc:
0.6664 - val_roc_auc: 0.1953 - learning_rate: 0.0010
Epoch 6/80
2/2          0s 26ms/step - loss:
0.7321 - pr_auc: 0.7344 - roc_auc: 0.5616 - val_loss: 0.9480 - val_pr_auc:
```

```
0.6619 - val_roc_auc: 0.2031 - learning_rate: 0.0010
Epoch 7/80
2/2          0s 24ms/step - loss:
0.7391 - pr_auc: 0.6068 - roc_auc: 0.5009 - val_loss: 0.9312 - val_pr_auc:
0.6619 - val_roc_auc: 0.2031 - learning_rate: 0.0010
Epoch 8/80
2/2          0s 26ms/step - loss:
0.7413 - pr_auc: 0.7443 - roc_auc: 0.5321 - val_loss: 0.9168 - val_pr_auc:
0.6712 - val_roc_auc: 0.2344 - learning_rate: 0.0010
Epoch 9/80
2/2          0s 26ms/step - loss:
0.7532 - pr_auc: 0.6878 - roc_auc: 0.4991 - val_loss: 0.9041 - val_pr_auc:
0.6735 - val_roc_auc: 0.2500 - learning_rate: 0.0010
Epoch 10/80
2/2          0s 25ms/step - loss:
0.7399 - pr_auc: 0.6602 - roc_auc: 0.5052 - val_loss: 0.8914 - val_pr_auc:
0.6808 - val_roc_auc: 0.2656 - learning_rate: 0.0010
Epoch 11/80
2/2          0s 36ms/step - loss:
0.7341 - pr_auc: 0.7087 - roc_auc: 0.4991 - val_loss: 0.8798 - val_pr_auc:
0.6761 - val_roc_auc: 0.2578 - learning_rate: 0.0010
Epoch 12/80
2/2          0s 39ms/step - loss:
0.7170 - pr_auc: 0.6671 - roc_auc: 0.5590 - val_loss: 0.8691 - val_pr_auc:
0.6761 - val_roc_auc: 0.2578 - learning_rate: 0.0010
Epoch 13/80
2/2          0s 25ms/step - loss:
0.6966 - pr_auc: 0.7723 - roc_auc: 0.6259 - val_loss: 0.8595 - val_pr_auc:
0.6761 - val_roc_auc: 0.2578 - learning_rate: 0.0010
Epoch 14/80
2/2          0s 24ms/step - loss:
0.7156 - pr_auc: 0.7130 - roc_auc: 0.5712 - val_loss: 0.8502 - val_pr_auc:
0.6761 - val_roc_auc: 0.2578 - learning_rate: 0.0010
Epoch 15/80
2/2          0s 24ms/step - loss:
0.7537 - pr_auc: 0.5546 - roc_auc: 0.3559 - val_loss: 0.8406 - val_pr_auc:
0.6808 - val_roc_auc: 0.2656 - learning_rate: 0.0010
Epoch 16/80
2/2          0s 24ms/step - loss:
0.7127 - pr_auc: 0.7145 - roc_auc: 0.5460 - val_loss: 0.8361 - val_pr_auc:
0.6752 - val_roc_auc: 0.2578 - learning_rate: 5.0000e-04
Epoch 17/80
2/2          0s 24ms/step - loss:
0.7234 - pr_auc: 0.5911 - roc_auc: 0.4618 - val_loss: 0.8318 - val_pr_auc:
0.6832 - val_roc_auc: 0.2734 - learning_rate: 5.0000e-04
Epoch 18/80
2/2          0s 24ms/step - loss:
0.7101 - pr_auc: 0.7010 - roc_auc: 0.5243 - val_loss: 0.8276 - val_pr_auc:
```

```
0.6867 - val_roc_auc: 0.2734 - learning_rate: 5.0000e-04
Epoch 19/80
2/2          0s 23ms/step - loss:
0.7148 - pr_auc: 0.6038 - roc_auc: 0.4800 - val_loss: 0.8238 - val_pr_auc:
0.6891 - val_roc_auc: 0.2812 - learning_rate: 5.0000e-04
Epoch 20/80
2/2          0s 25ms/step - loss:
0.7209 - pr_auc: 0.7052 - roc_auc: 0.5252 - val_loss: 0.8198 - val_pr_auc:
0.6832 - val_roc_auc: 0.2734 - learning_rate: 5.0000e-04
Epoch 21/80
2/2          0s 24ms/step - loss:
0.7123 - pr_auc: 0.7049 - roc_auc: 0.5252 - val_loss: 0.8161 - val_pr_auc:
0.6826 - val_roc_auc: 0.2734 - learning_rate: 5.0000e-04
Epoch 22/80
2/2          0s 27ms/step - loss:
0.7319 - pr_auc: 0.5634 - roc_auc: 0.4236 - val_loss: 0.8126 - val_pr_auc:
0.6891 - val_roc_auc: 0.2812 - learning_rate: 5.0000e-04
Epoch 23/80
2/2          0s 25ms/step - loss:
0.7113 - pr_auc: 0.6909 - roc_auc: 0.5122 - val_loss: 0.8093 - val_pr_auc:
0.6891 - val_roc_auc: 0.2812 - learning_rate: 5.0000e-04
Epoch 24/80
2/2          0s 25ms/step - loss:
0.7137 - pr_auc: 0.6952 - roc_auc: 0.5417 - val_loss: 0.8059 - val_pr_auc:
0.6906 - val_roc_auc: 0.2812 - learning_rate: 5.0000e-04
Epoch 25/80
2/2          0s 24ms/step - loss:
0.7061 - pr_auc: 0.6027 - roc_auc: 0.5009 - val_loss: 0.8027 - val_pr_auc:
0.6826 - val_roc_auc: 0.2734 - learning_rate: 5.0000e-04
Epoch 26/80
2/2          0s 24ms/step - loss:
0.6742 - pr_auc: 0.7485 - roc_auc: 0.6432 - val_loss: 0.7997 - val_pr_auc:
0.6826 - val_roc_auc: 0.2734 - learning_rate: 5.0000e-04
Epoch 27/80
2/2          0s 23ms/step - loss:
0.6750 - pr_auc: 0.7267 - roc_auc: 0.6615 - val_loss: 0.7969 - val_pr_auc:
0.6891 - val_roc_auc: 0.2891 - learning_rate: 5.0000e-04
Epoch 28/80
2/2          0s 24ms/step - loss:
0.7107 - pr_auc: 0.6950 - roc_auc: 0.5538 - val_loss: 0.7944 - val_pr_auc:
0.6891 - val_roc_auc: 0.2812 - learning_rate: 5.0000e-04
Epoch 29/80
2/2          0s 24ms/step - loss:
0.7280 - pr_auc: 0.5599 - roc_auc: 0.4089 - val_loss: 0.7918 - val_pr_auc:
0.6873 - val_roc_auc: 0.2734 - learning_rate: 5.0000e-04
Epoch 30/80
2/2          0s 26ms/step - loss:
0.6869 - pr_auc: 0.7796 - roc_auc: 0.6458 - val_loss: 0.7906 - val_pr_auc:
```

```
0.6953 - val_roc_auc: 0.2891 - learning_rate: 2.5000e-04
Epoch 31/80
2/2          0s 28ms/step - loss:
0.6950 - pr_auc: 0.6698 - roc_auc: 0.5868 - val_loss: 0.7896 - val_pr_auc:
0.6995 - val_roc_auc: 0.2969 - learning_rate: 2.5000e-04
Epoch 32/80
2/2          0s 44ms/step - loss:
0.6944 - pr_auc: 0.7330 - roc_auc: 0.5833 - val_loss: 0.7886 - val_pr_auc:
0.7019 - val_roc_auc: 0.3047 - learning_rate: 2.5000e-04
Epoch 33/80
2/2          0s 25ms/step - loss:
0.6839 - pr_auc: 0.6841 - roc_auc: 0.6102 - val_loss: 0.7873 - val_pr_auc:
0.7019 - val_roc_auc: 0.3047 - learning_rate: 2.5000e-04
Epoch 34/80
2/2          0s 24ms/step - loss:
0.7029 - pr_auc: 0.6720 - roc_auc: 0.5755 - val_loss: 0.7860 - val_pr_auc:
0.7001 - val_roc_auc: 0.2969 - learning_rate: 2.5000e-04
Epoch 35/80
2/2          0s 24ms/step - loss:
0.6913 - pr_auc: 0.7798 - roc_auc: 0.6189 - val_loss: 0.7848 - val_pr_auc:
0.7122 - val_roc_auc: 0.3047 - learning_rate: 2.5000e-04
Epoch 36/80
2/2          0s 24ms/step - loss:
0.6817 - pr_auc: 0.7860 - roc_auc: 0.6198 - val_loss: 0.7837 - val_pr_auc:
0.6808 - val_roc_auc: 0.2656 - learning_rate: 2.5000e-04
Epoch 37/80
2/2          0s 24ms/step - loss:
0.6762 - pr_auc: 0.8324 - roc_auc: 0.6901 - val_loss: 0.7824 - val_pr_auc:
0.6826 - val_roc_auc: 0.2734 - learning_rate: 2.5000e-04
Epoch 38/80
2/2          0s 23ms/step - loss:
0.7133 - pr_auc: 0.5926 - roc_auc: 0.4705 - val_loss: 0.7813 - val_pr_auc:
0.6854 - val_roc_auc: 0.2812 - learning_rate: 2.5000e-04
Epoch 39/80
2/2          0s 24ms/step - loss:
0.7141 - pr_auc: 0.6674 - roc_auc: 0.5200 - val_loss: 0.7802 - val_pr_auc:
0.6912 - val_roc_auc: 0.2891 - learning_rate: 2.5000e-04
Epoch 40/80
2/2          0s 25ms/step - loss:
0.6956 - pr_auc: 0.7603 - roc_auc: 0.6085 - val_loss: 0.7790 - val_pr_auc:
0.6937 - val_roc_auc: 0.2969 - learning_rate: 2.5000e-04
Epoch 41/80
2/2          0s 26ms/step - loss:
0.6917 - pr_auc: 0.6787 - roc_auc: 0.5842 - val_loss: 0.7783 - val_pr_auc:
0.6980 - val_roc_auc: 0.3047 - learning_rate: 1.2500e-04
Epoch 42/80
2/2          0s 24ms/step - loss:
0.6899 - pr_auc: 0.7789 - roc_auc: 0.6311 - val_loss: 0.7777 - val_pr_auc:
```

```
0.6980 - val_roc_auc: 0.3047 - learning_rate: 1.2500e-04
Epoch 43/80
2/2          0s 24ms/step - loss:
0.6884 - pr_auc: 0.7756 - roc_auc: 0.6380 - val_loss: 0.7771 - val_pr_auc:
0.7192 - val_roc_auc: 0.3281 - learning_rate: 1.2500e-04
Epoch 44/80
2/2          0s 24ms/step - loss:
0.6862 - pr_auc: 0.7753 - roc_auc: 0.6128 - val_loss: 0.7765 - val_pr_auc:
0.7192 - val_roc_auc: 0.3281 - learning_rate: 1.2500e-04
Epoch 45/80
2/2          0s 25ms/step - loss:
0.6870 - pr_auc: 0.7861 - roc_auc: 0.6128 - val_loss: 0.7758 - val_pr_auc:
0.7208 - val_roc_auc: 0.3359 - learning_rate: 1.2500e-04
Epoch 46/80
2/2          0s 23ms/step - loss:
0.7054 - pr_auc: 0.6354 - roc_auc: 0.4922 - val_loss: 0.7751 - val_pr_auc:
0.7156 - val_roc_auc: 0.3203 - learning_rate: 1.2500e-04
Epoch 47/80
2/2          0s 23ms/step - loss:
0.7021 - pr_auc: 0.7613 - roc_auc: 0.5729 - val_loss: 0.7745 - val_pr_auc:
0.7156 - val_roc_auc: 0.3203 - learning_rate: 1.2500e-04
Epoch 48/80
2/2          0s 24ms/step - loss:
0.6902 - pr_auc: 0.7521 - roc_auc: 0.6076 - val_loss: 0.7739 - val_pr_auc:
0.7156 - val_roc_auc: 0.3203 - learning_rate: 1.2500e-04
Epoch 49/80
2/2          0s 23ms/step - loss:
0.6653 - pr_auc: 0.7645 - roc_auc: 0.6693 - val_loss: 0.7733 - val_pr_auc:
0.7156 - val_roc_auc: 0.3203 - learning_rate: 1.2500e-04
Epoch 50/80
2/2          0s 24ms/step - loss:
0.6776 - pr_auc: 0.7635 - roc_auc: 0.6224 - val_loss: 0.7727 - val_pr_auc:
0.7172 - val_roc_auc: 0.3281 - learning_rate: 1.2500e-04
Epoch 51/80
2/2          0s 27ms/step - loss:
0.6874 - pr_auc: 0.6425 - roc_auc: 0.5521 - val_loss: 0.7725 - val_pr_auc:
0.7172 - val_roc_auc: 0.3281 - learning_rate: 6.2500e-05
Epoch 52/80
2/2          0s 27ms/step - loss:
0.6907 - pr_auc: 0.7084 - roc_auc: 0.5503 - val_loss: 0.7722 - val_pr_auc:
0.7172 - val_roc_auc: 0.3281 - learning_rate: 6.2500e-05
Epoch 53/80
2/2          0s 27ms/step - loss:
0.6591 - pr_auc: 0.8175 - roc_auc: 0.7196 - val_loss: 0.7720 - val_pr_auc:
0.7200 - val_roc_auc: 0.3359 - learning_rate: 6.2500e-05
Epoch 54/80
2/2          0s 31ms/step - loss:
0.6719 - pr_auc: 0.8261 - roc_auc: 0.6875 - val_loss: 0.7718 - val_pr_auc:
```

```

0.7200 - val_roc_auc: 0.3359 - learning_rate: 6.2500e-05
Epoch 55/80
2/2          0s 24ms/step - loss:
0.6867 - pr_auc: 0.7400 - roc_auc: 0.6250 - val_loss: 0.7715 - val_pr_auc:
0.7200 - val_roc_auc: 0.3359 - learning_rate: 6.2500e-05

```

1.0.9 9. Model B (sequence model): LSTM

We train an LSTM over the full sequence window. This can capture temporal patterns across tool-wear steps.

```
[10]: def build_lstm(timesteps: int, n_features: int):
    model = keras.Sequential([
        layers.Input(shape=(timesteps, n_features)),
        layers.LSTM(32, return_sequences=True),
        layers.Dropout(0.2),
        layers.LSTM(16),
        layers.Dropout(0.2),
        layers.Dense(1, activation="sigmoid"),
    ])
    model.compile(
        optimizer=keras.optimizers.Adam(learning_rate=1e-3),
        loss="binary_crossentropy",
        metrics=[keras.metrics.AUC(name="roc_auc"), keras.metrics.
        ↪AUC(curve="PR", name="pr_auc")]
    )
    return model

lstm = build_lstm(X_train_s.shape[1], X_train_s.shape[2])
lstm.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 10, 32)	4,736
dropout_2 (Dropout)	(None, 10, 32)	0
lstm_1 (LSTM)	(None, 16)	3,136
dropout_3 (Dropout)	(None, 16)	0
dense_3 (Dense)	(None, 1)	17

```
Total params: 7,889 (30.82 KB)
```

```
Trainable params: 7,889 (30.82 KB)
```

```
Non-trainable params: 0 (0.00 B)
```

```
[11]: history_lstm = lstm.fit(  
        X_train_s, y_train,  
        validation_data=(X_val_s, y_val),  
        epochs=80,  
        batch_size=32,  
        class_weight=class_weight_dict,  
        callbacks=callbacks,  
        verbose=1  
)
```

```
Epoch 1/80  
2/2          1s 234ms/step - loss:  
0.7000 - pr_auc: 0.5675 - roc_auc: 0.4401 - val_loss: 0.6682 - val_pr_auc:  
0.8459 - val_roc_auc: 0.5078 - learning_rate: 0.0010  
Epoch 2/80  
2/2          0s 29ms/step - loss:  
0.6992 - pr_auc: 0.5670 - roc_auc: 0.3819 - val_loss: 0.6753 - val_pr_auc:  
0.9064 - val_roc_auc: 0.6875 - learning_rate: 0.0010  
Epoch 3/80  
2/2          0s 28ms/step - loss:  
0.6988 - pr_auc: 0.5681 - roc_auc: 0.4427 - val_loss: 0.6827 - val_pr_auc:  
0.8930 - val_roc_auc: 0.6484 - learning_rate: 0.0010  
Epoch 4/80  
2/2          0s 30ms/step - loss:  
0.6972 - pr_auc: 0.5666 - roc_auc: 0.3759 - val_loss: 0.6890 - val_pr_auc:  
0.9139 - val_roc_auc: 0.7344 - learning_rate: 0.0010  
Epoch 5/80  
2/2          0s 27ms/step - loss:  
0.6915 - pr_auc: 0.6633 - roc_auc: 0.5165 - val_loss: 0.6925 - val_pr_auc:  
0.9050 - val_roc_auc: 0.6953 - learning_rate: 0.0010  
Epoch 6/80  
2/2          0s 28ms/step - loss:  
0.6889 - pr_auc: 0.6899 - roc_auc: 0.5408 - val_loss: 0.6982 - val_pr_auc:  
0.9044 - val_roc_auc: 0.7031 - learning_rate: 0.0010  
Epoch 7/80  
2/2          0s 28ms/step - loss:  
0.6890 - pr_auc: 0.7224 - roc_auc: 0.6155 - val_loss: 0.7040 - val_pr_auc:  
0.8976 - val_roc_auc: 0.6797 - learning_rate: 0.0010  
Epoch 8/80  
2/2          0s 27ms/step - loss:
```

```

0.6904 - pr_auc: 0.7430 - roc_auc: 0.5885 - val_loss: 0.7094 - val_pr_auc:
0.9001 - val_roc_auc: 0.6875 - learning_rate: 0.0010
Epoch 9/80
2/2          0s 29ms/step - loss:
0.6875 - pr_auc: 0.7033 - roc_auc: 0.6059 - val_loss: 0.7163 - val_pr_auc:
0.8924 - val_roc_auc: 0.6562 - learning_rate: 0.0010
Epoch 10/80
2/2          0s 28ms/step - loss:
0.6896 - pr_auc: 0.7042 - roc_auc: 0.5747 - val_loss: 0.7199 - val_pr_auc:
0.9053 - val_roc_auc: 0.6797 - learning_rate: 5.0000e-04
Epoch 11/80
2/2          0s 28ms/step - loss:
0.6908 - pr_auc: 0.6559 - roc_auc: 0.5964 - val_loss: 0.7231 - val_pr_auc:
0.9002 - val_roc_auc: 0.6797 - learning_rate: 5.0000e-04
Epoch 12/80
2/2          0s 27ms/step - loss:
0.6901 - pr_auc: 0.6809 - roc_auc: 0.5677 - val_loss: 0.7265 - val_pr_auc:
0.9029 - val_roc_auc: 0.6719 - learning_rate: 5.0000e-04
Epoch 13/80
2/2          0s 26ms/step - loss:
0.6854 - pr_auc: 0.6828 - roc_auc: 0.6345 - val_loss: 0.7302 - val_pr_auc:
0.8983 - val_roc_auc: 0.6641 - learning_rate: 5.0000e-04
Epoch 14/80
2/2          0s 26ms/step - loss:
0.6866 - pr_auc: 0.7525 - roc_auc: 0.6372 - val_loss: 0.7332 - val_pr_auc:
0.8911 - val_roc_auc: 0.6562 - learning_rate: 5.0000e-04

```

1.0.10 10. Evaluate models

We evaluate both models on the held-out test set using ROC-AUC and PR-AUC, plus classification report and confusion matrix.

```
[12]: def evaluate_classifier(name, y_true, y_prob, threshold=0.5):
    y_pred = (y_prob >= threshold).astype(int)
    print(name)
    print("ROC-AUC:", round(roc_auc_score(y_true, y_prob), 4))
    print("PR-AUC : ", round(average_precision_score(y_true, y_prob), 4))
    print(classification_report(y_true, y_pred, digits=4))
    print("Confusion matrix:\n", confusion_matrix(y_true, y_pred))
    print("-"*80)

proba_mlp = mlp.predict(X_test_last).ravel()
proba_lstm = lstm.predict(X_test_s).ravel()

evaluate_classifier("MLP (last timestep)", y_test, proba_mlp)
evaluate_classifier("LSTM (sequence)", y_test, proba_lstm)
```

```

6/6          0s 3ms/step
6/6          0s 20ms/step
MLP (last timestep)
ROC-AUC: 0.6075
PR-AUC : 0.7977

      precision    recall   f1-score   support

```

	precision	recall	f1-score	support
0	0.3267	0.9800	0.4900	50
1	0.9412	0.1368	0.2388	117
accuracy			0.3892	167
macro avg	0.6339	0.5584	0.3644	167
weighted avg	0.7572	0.3892	0.3140	167

Confusion matrix:

```

[[ 49   1]
 [101  16]]
-----
```

LSTM (sequence)

ROC-AUC: 0.5173
PR-AUC : 0.7387

	precision	recall	f1-score	support
0	0.3056	0.2200	0.2558	50
1	0.7023	0.7863	0.7419	117
accuracy			0.6168	167
macro avg	0.5039	0.5032	0.4989	167
weighted avg	0.5835	0.6168	0.5964	167

Confusion matrix:

```

[[11 39]
 [25 92]]
-----
```

1.0.11 11. ROC and Precision-Recall curves (2 columns)

We plot ROC (left) and Precision-Recall (right) for both deep models.

```
[13]: fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# ROC (left)
RocCurveDisplay.from_predictions(y_test, proba_mlp, name="MLP", ax=axes[0])
RocCurveDisplay.from_predictions(y_test, proba_lstm, name="LSTM", ax=axes[0])
axes[0].set_title("ROC Curve")

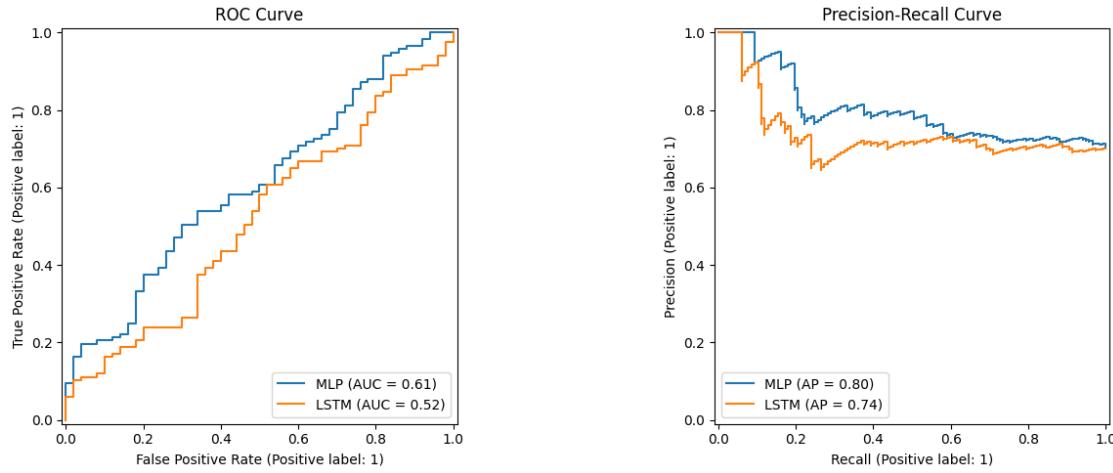
# PR (right)
```

```

PrecisionRecallDisplay.from_predictions(y_test, proba_mlp, name="MLP", ax=axes[1])
PrecisionRecallDisplay.from_predictions(y_test, proba_lstm, name="LSTM", ax=axes[1])
axes[1].set_title("Precision-Recall Curve")

plt.tight_layout()
plt.show()

```



1.0.12 12. Training curves

We visualize learning curves (loss and PR-AUC) to check convergence and overfitting.

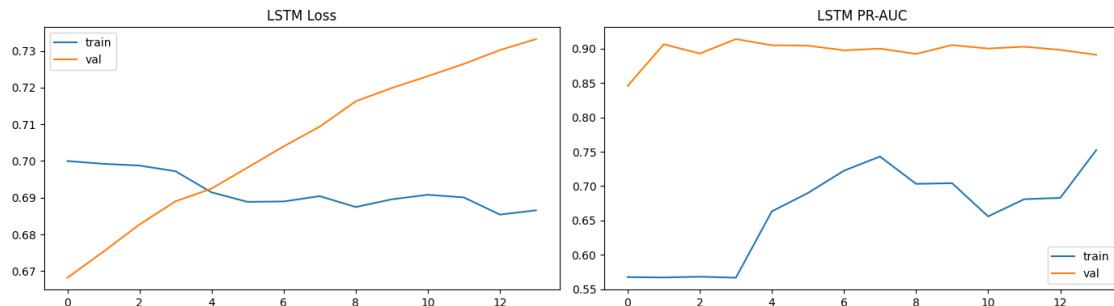
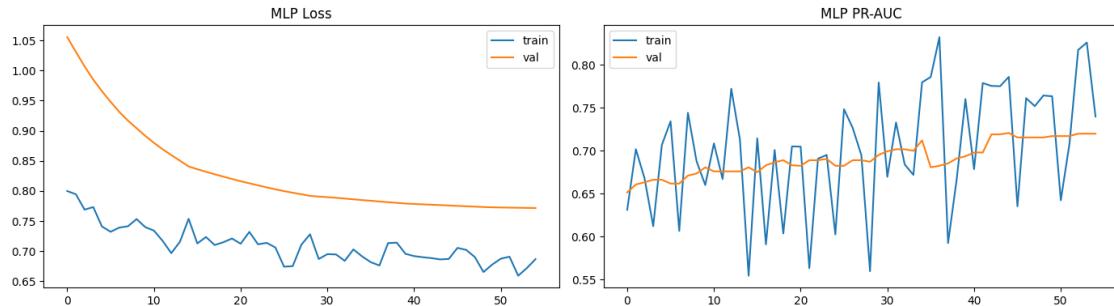
```
[14]: def plot_history(hist, title_prefix):
    hist_df = pd.DataFrame(hist.history)
    fig, axes = plt.subplots(1, 2, figsize=(14, 4))

    axes[0].plot(hist_df["loss"], label="train")
    axes[0].plot(hist_df["val_loss"], label="val")
    axes[0].set_title(f"{title_prefix} Loss")
    axes[0].legend()

    if "pr_auc" in hist_df.columns:
        axes[1].plot(hist_df["pr_auc"], label="train")
        axes[1].plot(hist_df["val_pr_auc"], label="val")
        axes[1].set_title(f"{title_prefix} PR-AUC")
        axes[1].legend()

    plt.tight_layout()
    plt.show()
```

```
plot_history(history_mlp, "MLP")
plot_history(history_lstm, "LSTM")
```



1.0.13 13. Save models

This saves the trained Keras models for reproducibility.

```
[15]: OUT_DIR = "../models"
os.makedirs(OUT_DIR, exist_ok=True)

mlp_path = os.path.join(OUT_DIR, "model2_mlp.keras")
lstm_path = os.path.join(OUT_DIR, "model2_lstm.keras")

mlp.save(mlp_path)
lstm.save(lstm_path)

print("Saved:", mlp_path)
print("Saved:", lstm_path)
```

```
Saved: ../models/model2_mlp.keras
Saved: ../models/model2_lstm.keras
```

1.0.14 14. Export Deep Learning Results for Tableau

```
[20]: os.makedirs("../outputs", exist_ok=True)

# Required variables from this notebook
required_vars = ["df_tw", "TIME_COL", "SEQ_LEN", "val_end", "y_test"]
missing = [v for v in required_vars if v not in globals()]
if missing:
    raise NameError(f"Missing required variables for export: {missing}. Run the"
                   "earlier cells first.")

# Build tool_wear aligned with sequence samples
# X_seq / y_seq are created from df_tw starting at index SEQ_LEN-1 (label at
↳ window end)
tool_wear_seq = df_tw[TIME_COL].iloc[SEQ_LEN - 1:].values # length ==
    ↳ len(y_seq)

# Test slice corresponds to X_seq[val_end:], y_seq[val_end:]
tool_wear_test = tool_wear_seq[val_end:]

# Get probabilities
if "proba_lstm" in globals():
    lstm_probs = np.asarray(proba_lstm).ravel()
else:
    # fallback: compute from model + scaled sequences if needed
    if "lstm" not in globals() or "X_test_s" not in globals():
        raise NameError("Missing proba_lstm and cannot recompute (need lstm and"
                       "↳ X_test_s).")
    lstm_probs = np.asarray(lstm.predict(X_test_s)).ravel()

if "proba_mlp" in globals():
    mlp_probs = np.asarray(proba_mlp).ravel()
else:
    if "mlp" not in globals() or "X_test_last" not in globals():
        raise NameError("Missing proba_mlp and cannot recompute (need mlp and"
                       "↳ X_test_last).")
    mlp_probs = np.asarray(mlp.predict(X_test_last)).ravel()

# Safety: ensure lengths match
n = len(y_test)
if len(tool_wear_test) != n:
    raise ValueError(f"tool_wear_test length ({len(tool_wear_test)}) != y_test"
                     "length ({n}).")
if len(lstm_probs) != n:
    raise ValueError(f"lstm_probs length ({len(lstm_probs)}) != y_test length"
                     "({n}).")
if len(mlp_probs) != n:
```

```

    raise ValueError(f"mlp_probs length ({len(mlp_probs)}) != y_test length_{n}.")
# Recommended threshold (Youden's J) using LSTM
fpr, tpr, thresholds = roc_curve(y_test, lstm_probs)
optimal_idx = (tpr - fpr).argmax()
optimal_threshold = float(thresholds[optimal_idx])
print(f"Recommended threshold (LSTM - Youden's J): {optimal_threshold:.6f}")

# Predicted labels for Tableau filters
pred_label_lstm = (lstm_probs >= optimal_threshold).astype(int)
pred_label_mlp = (mlp_probs >= optimal_threshold).astype(int)

# Export
export_df = pd.DataFrame({
    "tool_wear": tool_wear_test,
    "failure_actual": np.asarray(y_test).astype(int),
    "failure_prob_lstm": lstm_probs,
    "failure_prob_mlp": mlp_probs,
    "pred_label_lstm": pred_label_lstm,
    "pred_label_mlp": pred_label_mlp,
    "recommended_threshold": optimal_threshold
}).sort_values("tool_wear")

export_path = "../outputs/pred_failure_deep.csv"
export_df.to_csv(export_path, index=False)

print(f"Saved: {export_path} | rows={len(export_df)}")

```

Recommended threshold (LSTM - Youden's J): 0.503694

Saved: ../outputs/pred_failure_deep.csv | rows=167

1.0.15 14. Summary

In this notebook, we implemented deep learning models for predictive maintenance using a time-aware formulation of the AI4I dataset.

Key points:

- Tool wear [min] was used as a proxy for temporal progression.
- Data were aggregated by tool wear, preserving the binary failure label using `max()`.
- Sliding-window sequences were created to enable temporal modeling.
- Two deep models were trained:
 - MLP (last timestep baseline)
 - LSTM (sequence model)

Results show that:

- The MLP serves as a strong tabular deep-learning baseline.
- The LSTM captures temporal dependencies across tool-wear progression.

- Precision-Recall AUC is particularly important due to class imbalance.

Overall, Model 2 demonstrates how sequence-based deep learning can extend traditional machine learning approaches for failure prediction, while maintaining strict time-aware splitting to avoid data leakage.